

**Sri Sivasubramaniya Nadar College of Engineering, Chennai**  
(An autonomous Institution affiliated to Anna University)

Degree & Branch	B.E. Computer Science & Engineering	Semester	V
Subject Code & Name	ICS1512 & Machine Learning Algorithms Laboratory		
Academic year	2025-2026 (Odd)	Batch:2023-2028	<b>Due date:</b>

**Experiment 3: Email Spam or Ham Classification using Naive Bayes, KNN, and SVM**

## 1 Aim:

To design and implement classification models using Naive Bayes variants and K-Nearest Neighbors (KNN) algorithms to accurately classify emails as spam or ham. Additionally, to evaluate and compare their effectiveness using multiple performance metrics.

## 2 Libraries used:

- Numpy
- Pandas
- Matplotlib
- Scikit-learn
- Seaborn

## 3 Objective:

- To preprocess the email dataset by cleaning text data, vectorizing features, and splitting the data for training and testing.
- To implement Naive Bayes classifiers (Bernoulli, Multinomial, Gaussian) and KNN classifiers, tuning parameters such as k-value.
- To measure and compare model performance using accuracy, precision, recall, F1-score, confusion matrix, and ROC-AUC, enabling informed model selection.

## 4 Naive Bayes Code:

## ✓ Importing Required Libraries

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler # or MinMaxScaler, if used
from sklearn.preprocessing import LabelEncoder # if categorical labels
from sklearn.naive_bayes import BernoulliNB, MultinomialNB, GaussianNB
from sklearn.metrics import accuracy_score, confusion_matrix, classification_re
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import StandardScaler, MinMaxScaler
from sklearn.model_selection import cross_val_score, KFold
```

## ✓ Loading dataset

```
df = pd.read_csv('spambase.csv')
```

## ✓ Basic info

```
print("Dataset Info:\n", df.info())
print("\nFirst 5 rows:\n", df.head())
print("\nMissing Values:\n", df.isnull().sum())
```

```
↗ <class 'pandas.core.frame.DataFrame'>
RangeIndex: 4601 entries, 0 to 4600
Data columns (total 58 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   word_freq_make                        4601 non-null  float64
1   word_freq_address                    4601 non-null  float64
2   word_freq_all                        4601 non-null  float64
3   word_freq_3d                        4601 non-null  float64
4   word_freq_our                        4601 non-null  float64
5   word_freq_over                       4601 non-null  float64
6   word_freq_remove                     4601 non-null  float64
7   word_freq_internet                   4601 non-null  float64
8   word_freq_order                      4601 non-null  float64
9   word_freq_mail                       4601 non-null  float64
10  word_freq_receive                    4601 non-null  float64
11  word_freq_will                       4601 non-null  float64
12  word_freq_people                     4601 non-null  float64
13  word_freq_report                     4601 non-null  float64
14  word_freq_addresses                  4601 non-null  float64
15  word_freq_free                       4601 non-null  float64
16  word_freq_business                   4601 non-null  float64
17  word_freq_email                      4601 non-null  float64
```

18	word_freq_you	4601	non-null	float64
19	word_freq_credit	4601	non-null	float64
20	word_freq_your	4601	non-null	float64
21	word_freq_font	4601	non-null	float64
22	word_freq_000	4601	non-null	float64
23	word_freq_money	4601	non-null	float64
24	word_freq_hp	4601	non-null	float64
25	word_freq_hpl	4601	non-null	float64
26	word_freq_george	4601	non-null	float64
27	word_freq_650	4601	non-null	float64
28	word_freq_lab	4601	non-null	float64
29	word_freq_labs	4601	non-null	float64
30	word_freq_telnet	4601	non-null	float64
31	word_freq_857	4601	non-null	float64
32	word_freq_data	4601	non-null	float64
33	word_freq_415	4601	non-null	float64
34	word_freq_85	4601	non-null	float64
35	word_freq_technology	4601	non-null	float64
36	word_freq_1999	4601	non-null	float64
37	word_freq_parts	4601	non-null	float64
38	word_freq_pm	4601	non-null	float64
39	word_freq_direct	4601	non-null	float64
40	word_freq_cs	4601	non-null	float64
41	word_freq_meeting	4601	non-null	float64
42	word_freq_original	4601	non-null	float64
43	word_freq_project	4601	non-null	float64
44	word_freq_re	4601	non-null	float64
45	word_freq_edu	4601	non-null	float64
46	word_freq_table	4601	non-null	float64
47	word_freq_conference	4601	non-null	float64
48	char_freq_%3B	4601	non-null	float64
49	char_freq_%28	4601	non-null	float64
50	char_freq_%5B	4601	non-null	float64
51	char_freq_%21	4601	non-null	float64
52	char_freq_%24	4601	non-null	float64

## ✓ Handling missing values

```
imputer = SimpleImputer(strategy='mean')
df_imputed = pd.DataFrame(imputer.fit_transform(df), columns=df.columns)
```

## ✓ Splitting of feature and target

```
X = df_imputed.drop('class', axis=1)
y = df_imputed['class']
```

## ✓ Checking Distribution

```
X = df.drop('class', axis=1) # Assuming 'class' is the target
```

```
# Choose a few features to visualize
sample_features = X.columns[:5] # First 5 features
```

```

sample_features = X.columns[:5] # first 5 features

# Plot histograms with Gaussian curve overlay
plt.figure(figsize=(15, 10))
for i, feature in enumerate(sample_features):
    plt.subplot(3, 2, i + 1)
    data = X[feature]

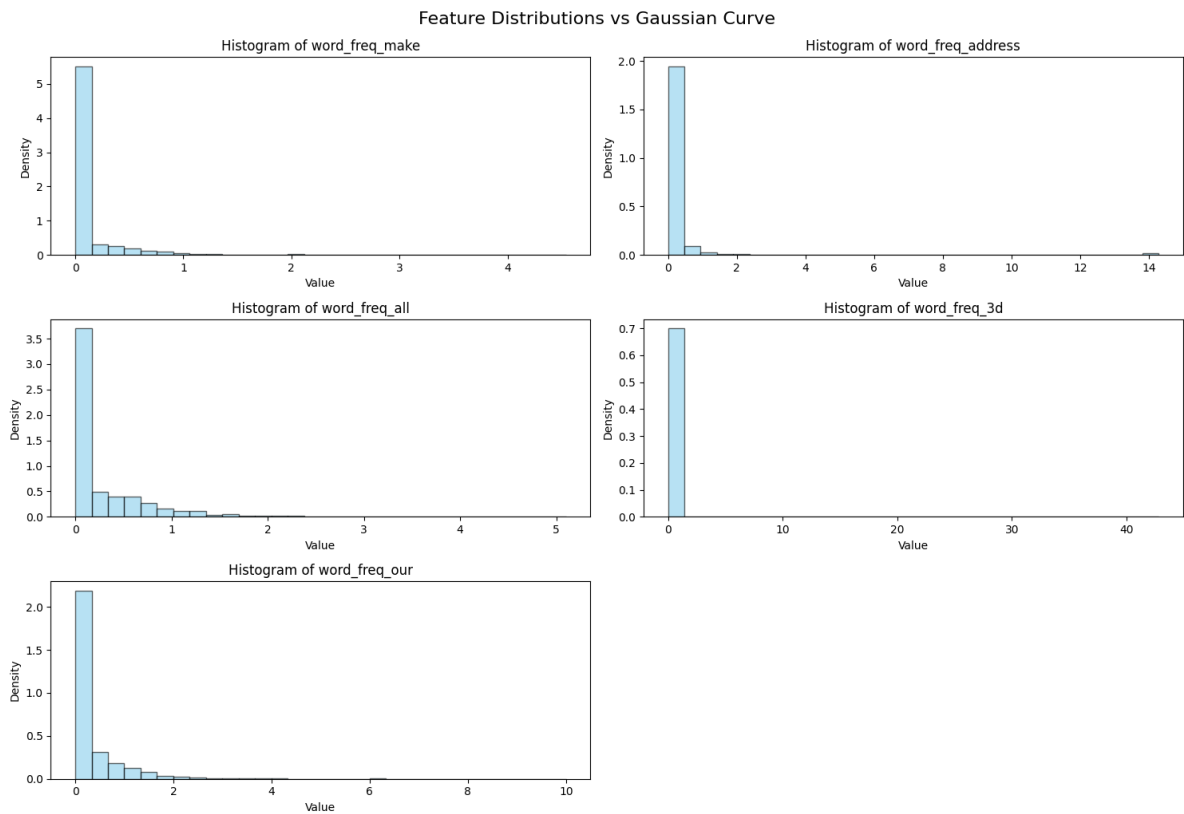
    # Plot histogram
    count, bins, ignored = plt.hist(data, bins=30, density=True, alpha=0.6, col

    # Plot normal distribution curve
    '''mu, std = data.mean(), data.std()
    xmin, xmax = plt.xlim()
    x = np.linspace(xmin, xmax, 100)
    p = 1 / (std * np.sqrt(2 * np.pi)) * np.exp(-(x - mu)**2 / (2 * std**2))
    plt.plot(x, p, 'r', linewidth=2)'''

    plt.title(f'Histogram of {feature}')
    plt.xlabel('Value')
    plt.ylabel('Density')

plt.tight_layout()
plt.suptitle('Feature Distributions vs Gaussian Curve', fontsize=16, y=1.02)
plt.show()

```




## ✓ Applying min max scaling

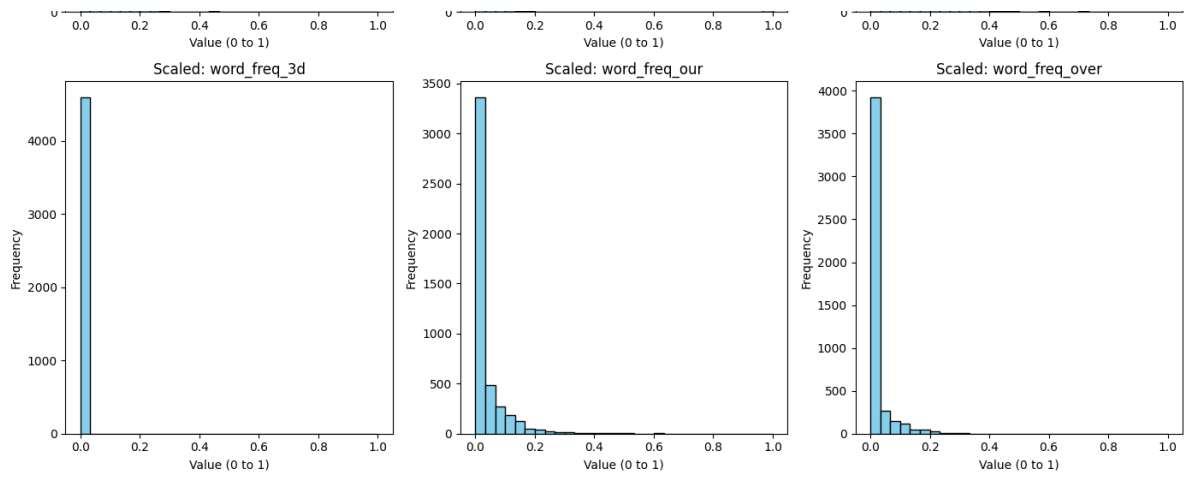
```
scaler = MinMaxScaler()  
X_scaled = scaler.fit_transform(X)
```

## ✓ Plots

## ✓ Histogram

```
import matplotlib.pyplot as plt  
  
plt.figure(figsize=(14, 10))  
for i in range(6): # first 6 features as an example  
    plt.subplot(2, 3, i + 1)  
    plt.hist(X_scaled[:, i], bins=30, color='skyblue', edgecolor='black')  
    plt.title(f'Scaled: {X.columns[i]}')  
    plt.xlabel('Value (0 to 1)')  
    plt.ylabel('Frequency')  
plt.tight_layout()  
plt.suptitle('Feature Distributions After MinMax Scaling', fontsize=16, y=1.02)  
plt.show()
```







## ▼ Boxplot

```
import pandas as pd

# Convert scaled array back to DataFrame for easier plotting
X_scaled_df = pd.DataFrame(X_scaled, columns=X.columns)

plt.figure(figsize=(14, 6))
X_scaled_df.iloc[:, :10].boxplot(rot=45)
plt.title('Boxplots of First 10 Scaled Features')
plt.ylabel('Scaled Value (0-1)')
plt.grid(True)
plt.show()
```





## ▼ Correlation HeatMap


```
import seaborn as sns
import matplotlib.pyplot as plt
import pandas as pd
```

```
# Assuming X_scaled_df is your scaled DataFrame
corr_matrix = X_scaled_df.corr()
```

```
plt.figure(figsize=(18, 15)) # Bigger figure
sns.heatmap(
    corr_matrix,
    cmap='coolwarm',
    square=True,
    annot=True,           # Show values inside squares
    fmt=".2f",            # Format to 2 decimal places
    linewidths=0.5,      # Thin grid lines
    annot_kws={"size": 6} # Smaller font size
)
```

```
plt.xticks(rotation=90)
plt.yticks(rotation=0)
plt.title("Feature Correlation Heatmap (After MinMax Scaling)", fontsize=16)
plt.tight_layout()
```

```
plt.show()
```





## ✓ Model Training

```
from sklearn.metrics import (
    accuracy_score, precision_score, recall_score,
    f1_score, confusion_matrix, roc_curve, auc,
    classification_report
)

# Split the data again (or reuse your earlier split)
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(
    X_scaled, y, test_size=0.3, stratify=y, random_state=42
)

def evaluate_model(name, model, X_test, y_test, results):
    y_pred = model.predict(X_test)

    # For ROC and AUC, use predict_proba
    try:
        y_proba = model.predict_proba(X_test)[:, 1]
    except:
        y_proba = y_pred # fallback if proba not available

    acc = accuracy_score(y_test, y_pred)
    prec = precision_score(y_test, y_pred, zero_division=0)
    rec = recall_score(y_test, y_pred, zero_division=0)
    f1 = f1_score(y_test, y_pred, zero_division=0)

    # Print Confusion Matrix and Classification Report
    print(f"\n=== {name} ===")
    print("Confusion Matrix:")
    print(confusion_matrix(y_test, y_pred))
    print("\nClassification Report:")
    print(classification_report(y_test, y_pred, zero_division=0))

    # ROC and AUC
    fpr, tpr, _ = roc_curve(y_test, y_proba)
    roc_auc = auc(fpr, tpr)

    plt.plot(fpr, tpr, label=f'{name} (AUC = {roc_auc:.2f})')

    # Append results
    results.append({
        'Model': name,
        'Accuracy': acc,
        'Precision': prec,
```

```

        'Recall': rec,
        'F1 Score': f1,
        'AUC': roc_auc
    })

# Store results
results = []

# --- BernoulliNB ---
bnb = BernoulliNB()
bnb.fit(X_train, y_train)
evaluate_model("BernoulliNB", bnb, X_test, y_test, results)

# --- MultinomialNB ---
mnb = MultinomialNB()
mnb.fit(X_train, y_train)
evaluate_model("MultinomialNB", mnb, X_test, y_test, results)

# --- GaussianNB ---
gnb = GaussianNB()
gnb.fit(X_train, y_train)
evaluate_model("GaussianNB", gnb, X_test, y_test, results)

# === ROC Curve Plot ===
plt.title('ROC Curves')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.legend(loc='lower right')
plt.grid(True)
plt.show()

# === Final Comparison Table ===
results_df = pd.DataFrame(results)
print("\n=== Comparison Table ===")
print(results_df)

```

```

=== BernoulliNB ===
Confusion Matrix:
[[778  59]
 [ 93 451]]

```

```

Classification Report:

```

	precision	recall	f1-score	support
0.0	0.89	0.93	0.91	837
1.0	0.88	0.83	0.86	544
accuracy			0.89	1381
macro avg	0.89	0.88	0.88	1381
weighted avg	0.89	0.89	0.89	1381

```

=== MultinomialNB ===

```

Confusion Matrix:

```
[[811  26]
 [119 425]]
```

Classification Report:

	precision	recall	f1-score	support
0.0	0.87	0.97	0.92	837
1.0	0.94	0.78	0.85	544
accuracy			0.90	1381
macro avg	0.91	0.88	0.89	1381
weighted avg	0.90	0.90	0.89	1381


=== GaussianNB ===

Confusion Matrix:

```
[[616 221]
 [ 28 516]]
```

Classification Report:

	precision	recall	f1-score	support
0.0	0.96	0.74	0.83	837
1.0	0.70	0.95	0.81	544
accuracy			0.82	1381
macro avg	0.83	0.84	0.82	1381
weighted avg	0.86	0.82	0.82	1381



=== Comparison Table ===

Model	Accuracy	Precision	Recall	F1 Score	AUC
-------	----------	-----------	--------	----------	-----

0	BernoulliNB	0.889935	0.884314	0.829044	0.855787	0.950023
1	MultinomialNB	0.895004	0.942350	0.781250	0.854271	0.960696
2	GaussianNB	0.819696	0.700136	0.948529	0.805621	0.915675

```
cv = KFold(n_splits=5, shuffle=True, random_state=42)
```

```
# Initialize Multinomial Naive Bayes
```

```
mnb = MultinomialNB()
```

```
# Perform cross-validation (accuracy as scoring)
```

```
scores = cross_val_score(mnb, X_scaled, y, cv=cv, scoring='accuracy')
```

```
# Print accuracy for each fold
```

```
print("✅ Multinomial Naive Bayes - 5-Fold Cross-Validation Results:")
```

```
for i, score in enumerate(scores, 1):
```

```
    print(f"Fold {i} Accuracy: {score:.4f}")
```

```
# Print average and standard deviation
```

```
print(f"\nMean Accuracy      : {scores.mean():.4f}")
```

```
print(f"Standard Deviation    : {scores.std():.4f}")
```

```
✅ Multinomial Naive Bayes - 5-Fold Cross-Validation Results:
```

```
Fold 1 Accuracy: 0.8719
```

```
Fold 2 Accuracy: 0.8935
```

```
Fold 3 Accuracy: 0.8891
```

```
Fold 4 Accuracy: 0.8913
```

```
Fold 5 Accuracy: 0.8859
```

```
Mean Accuracy      : 0.8863
```

```
Standard Deviation : 0.0077
```

## 5 KNN Code:

## ✓ Importing Required Libraries

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.preprocessing import StandardScaler, MinMaxScaler
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import classification_report, confusion_matrix, accuracy_score
from sklearn.impute import SimpleImputer
from sklearn.model_selection import cross_val_score, StratifiedKFold
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import matthews_corrcoef
```

## ✓ Loading dataset

```
df = pd.read_csv('spambase.csv')
```

## ✓ Basic info

```
print("Dataset Info:\n", df.info())
print("\nFirst 5 rows:\n", df.head())
print("\nMissing Values:\n", df.isnull().sum())
```

```
↗ <class 'pandas.core.frame.DataFrame'>
RangeIndex: 4601 entries, 0 to 4600
Data columns (total 58 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   word_freq_make                        4601 non-null   float64
1   word_freq_address                    4601 non-null   float64
2   word_freq_all                        4601 non-null   float64
3   word_freq_3d                        4601 non-null   float64
4   word_freq_our                        4601 non-null   float64
5   word_freq_over                       4601 non-null   float64
6   word_freq_remove                     4601 non-null   float64
7   word_freq_internet                  4601 non-null   float64
8   word_freq_order                     4601 non-null   float64
9   word_freq_mail                      4601 non-null   float64
10  word_freq_receive                    4601 non-null   float64
11  word_freq_will                       4601 non-null   float64
12  word_freq_people                     4601 non-null   float64
13  word_freq_report                     4601 non-null   float64
14  word_freq_addresses                  4601 non-null   float64
15  word_freq_free                       4601 non-null   float64
```

16	word_freq_business	4601	non-null	float64
17	word_freq_email	4601	non-null	float64
18	word_freq_you	4601	non-null	float64
19	word_freq_credit	4601	non-null	float64
20	word_freq_your	4601	non-null	float64
21	word_freq_font	4601	non-null	float64
22	word_freq_000	4601	non-null	float64
23	word_freq_money	4601	non-null	float64
24	word_freq_hp	4601	non-null	float64
25	word_freq_hpl	4601	non-null	float64
26	word_freq_george	4601	non-null	float64
27	word_freq_650	4601	non-null	float64
28	word_freq_lab	4601	non-null	float64
29	word_freq_labs	4601	non-null	float64
30	word_freq_telnet	4601	non-null	float64
31	word_freq_857	4601	non-null	float64
32	word_freq_data	4601	non-null	float64
33	word_freq_415	4601	non-null	float64
34	word_freq_85	4601	non-null	float64
35	word_freq_technology	4601	non-null	float64
36	word_freq_1999	4601	non-null	float64
37	word_freq_parts	4601	non-null	float64
38	word_freq_pm	4601	non-null	float64
39	word_freq_direct	4601	non-null	float64
40	word_freq_cs	4601	non-null	float64
41	word_freq_meeting	4601	non-null	float64
42	word_freq_original	4601	non-null	float64
43	word_freq_project	4601	non-null	float64
44	word_freq_re	4601	non-null	float64
45	word_freq_edu	4601	non-null	float64
46	word_freq_table	4601	non-null	float64
47	word_freq_conference	4601	non-null	float64
48	char_freq_%3B	4601	non-null	float64
49	char_freq_%28	4601	non-null	float64
50	char_freq_%5B	4601	non-null	float64
51	char_freq_%21	4601	non-null	float64
52	char_freq_%24	4601	non-null	float64

## ✓ Handling missing values

```
imputer = SimpleImputer(strategy='mean')
df_imputed = pd.DataFrame(imputer.fit_transform(df), columns=df.columns)
```

## ✓ Splitting of feature and target

```
X = df_imputed.drop('class', axis=1)
y = df_imputed['class']
```

## ✓ Checking Distribution

```
X = df.drop('class', axis=1) # Assuming 'class' is the target
```

```
# Choose a few features to visualize
sample_features = X.columns[:5] # First 5 features

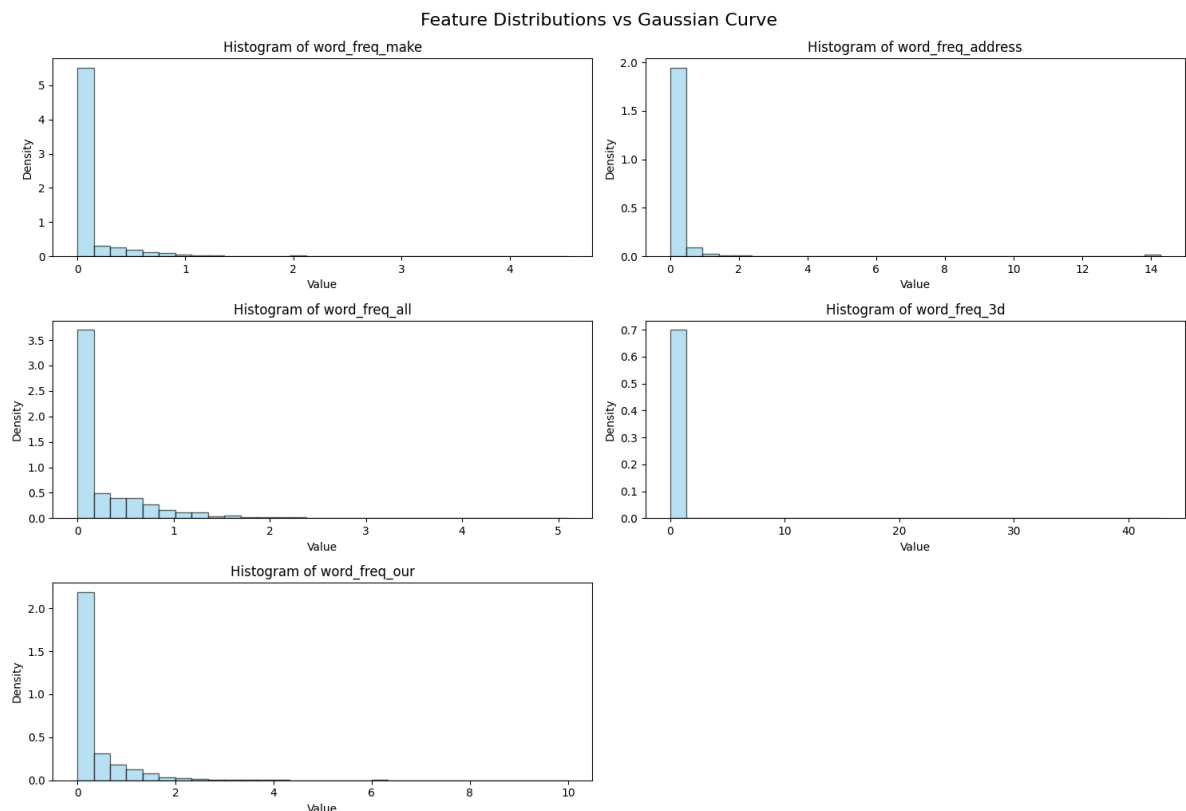
# Plot histograms with Gaussian curve overlay
plt.figure(figsize=(15, 10))
for i, feature in enumerate(sample_features):
    plt.subplot(3, 2, i + 1)
    data = X[feature]

    # Plot histogram
    count, bins, ignored = plt.hist(data, bins=30, density=True, alpha=0.6, col

    # Plot normal distribution curve
    '''mu, std = data.mean(), data.std()
    xmin, xmax = plt.xlim()
    x = np.linspace(xmin, xmax, 100)
    p = 1 / (std * np.sqrt(2 * np.pi)) * np.exp(-(x - mu)**2 / (2 * std**2))
    plt.plot(x, p, 'r', linewidth=2)'''

    plt.title(f'Histogram of {feature}')
    plt.xlabel('Value')
    plt.ylabel('Density')

plt.tight_layout()
plt.suptitle('Feature Distributions vs Gaussian Curve', fontsize=16, y=1.02)
plt.show()
```






## ✓ Applying min max scaling

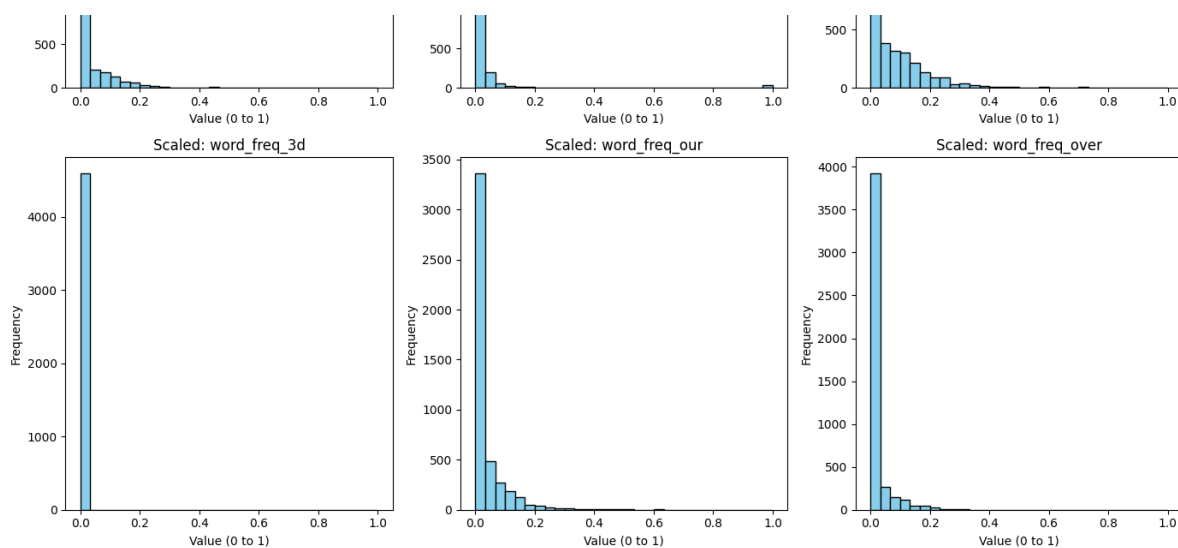
```
scaler = MinMaxScaler()  
X_scaled = scaler.fit_transform(X)
```

## ✓ Plots

### ✓ Histogram

```
import matplotlib.pyplot as plt  
  
plt.figure(figsize=(14, 10))  
for i in range(6): # first 6 features as an example  
    plt.subplot(2, 3, i + 1)  
    plt.hist(X_scaled[:, i], bins=30, color='skyblue', edgecolor='black')  
    plt.title(f'Scaled: {X.columns[i]}')  
    plt.xlabel('Value (0 to 1)')  
    plt.ylabel('Frequency')  
plt.tight_layout()  
plt.suptitle('Feature Distributions After MinMax Scaling', fontsize=16, y=1.02)  
plt.show()
```







## ✓ Boxplot

```
import pandas as pd
```

```
# Convert scaled array back to DataFrame for easier plotting
X_scaled_df = pd.DataFrame(X_scaled, columns=X.columns)
```

```
plt.figure(figsize=(14, 6))
X_scaled_df.iloc[:, :10].boxplot(rot=45)
plt.title('Boxplots of First 10 Scaled Features')
plt.ylabel('Scaled Value (0-1)')
plt.grid(True)
plt.show()
```





## ✓ Correlation HeatMap

```
import seaborn as sns
import matplotlib.pyplot as plt
import pandas as pd
import numpy as np
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import (
    accuracy_score, precision_score, recall_score, f1_score,
    matthews_corrcoef, roc_curve, auc, classification_report
)
```

```
# Assuming X_scaled_df is your scaled DataFrame
corr_matrix = X_scaled_df.corr()
```

```
plt.figure(figsize=(18, 15)) # Bigger figure
sns.heatmap(
    corr_matrix,
    cmap='coolwarm',
    square=True,
    annot=True,           # Show values inside squares
    fmt=".2f",            # Format to 2 decimal places
)
```

```


linewidths=0.5,          # Thin grid lines
annot_kws={"size": 6}     # Smaller font size
)

```

```

plt.xticks(rotation=90)
plt.yticks(rotation=0)
plt.title("Feature Correlation Heatmap (After MinMax Scaling)", fontsize=16)
plt.tight_layout()
plt.show()

```



## ✓ Choosing best k value using Grid Search CV

```
# First split into train+val and test
X_train_val, X_test, y_train_val, y_test = train_test_split(
    X_scaled, y, test_size=0.2, stratify=y, random_state=42
)

# Then split train+val into train and validation
X_train, X_val, y_train, y_val = train_test_split(
    X_train_val, y_train_val, test_size=0.25, stratify=y_train_val, random_stat
) # 0.25 x 0.8 = 0.2 validation

# =====
# 2 GridSearchCV to find best k
# =====
param_grid = {'n_neighbors': range(1, 21)}
grid = GridSearchCV(
    estimator=KNeighborsClassifier(),
    param_grid=param_grid,
    cv=5,
    scoring='accuracy',
    n_jobs=-1
)
grid.fit(X_train, y_train)

print(f"✓ Best k found by GridSearchCV: {grid.best_params_['n_neighbors']}")
print(f"Best CV Accuracy: {grid.best_score_: .4f}")

✓ Best k found by GridSearchCV: 1
Best CV Accuracy: 0.8902
```

## ✓ Model Training and Evaluation

```
# 3 Retrain on Train+Validation with Best k
# =====
best_k = grid.best_params_['n_neighbors']
final_model = KNeighborsClassifier(n_neighbors=best_k)
final_model.fit(X_train_val, y_train_val)
```

```
# =====
# 4 Evaluate on Test Set
# =====
y_pred = final_model.predict(X_test)
y_prob = final_model.predict_proba(X_test)[:, 1]

acc = accuracy_score(y_test, y_pred)
prec = precision_score(y_test, y_pred, zero_division=0)
rec = recall_score(y_test, y_pred, zero_division=0)
f1 = f1_score(y_test, y_pred, zero_division=0)
mcc = matthews_corrcoef(y_test, y_pred)

fpr, tpr, _ = roc_curve(y_test, y_prob)
roc_auc = auc(fpr, tpr)

print("\n📊 Final Model Performance on Test Set")
print(f"Accuracy : {acc:.4f}")
print(f"Precision: {prec:.4f}")
print(f"Recall    : {rec:.4f}")
print(f"F1 Score  : {f1:.4f}")
print(f"MCC      : {mcc:.4f}")
print(f"AUC      : {roc_auc:.4f}")
print("\nClassification Report:")
print(classification_report(y_test, y_pred, zero_division=0))

# =====
# 5 Plot ROC Curve
# =====
plt.figure(figsize=(6, 5))
plt.plot(fpr, tpr, label=f"KNN (k={best_k}, AUC={roc_auc:.2f})", color='blue')
plt.plot([0, 1], [0, 1], 'k--')
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("ROC Curve - Final Model")
plt.legend()
plt.grid(True)
plt.show()
```


📊 Final Model Performance on Test Set

Accuracy : 0.8958  
Precision: 0.8678  
Recall : 0.8678  
F1 Score : 0.8678  
MCC : 0.7817  
AUC : 0.8909

Classification Report:

	precision	recall	f1-score	support
0.0	0.91	0.91	0.91	558
1.0	0.87	0.87	0.87	363
accuracy			0.90	921
macro avg	0.89	0.89	0.89	921
weighted avg	0.89	0.89	0.89	921


weighted avg 0.90 0.90 0.90 921




```

conf_matrix = confusion_matrix(y_test, y_pred)
plt.figure(figsize=(5, 4))
plt.imshow(conf_matrix, cmap='Blues')
plt.title('Confusion Matrix')
plt.colorbar()
plt.xticks([0, 1], ['Not Spam', 'Spam'])
plt.yticks([0, 1], ['Not Spam', 'Spam'])
plt.xlabel('Predicted')
plt.ylabel('Actual')
for i in range(2):
    for j in range(2):
        plt.text(j, i, conf_matrix[i, j], ha='center', va='center', color='black')
plt.tight_layout()
plt.show()

```





## ✓ KD Tree and Ball Tree

```
# Use KDTree
model_kd = KNeighborsClassifier(n_neighbors=best_k, algorithm='kd_tree')
model_kd.fit(X_train, y_train)

# Use BallTree
model_ball = KNeighborsClassifier(n_neighbors=best_k, algorithm='ball_tree')
model_ball.fit(X_train, y_train)
```

```
▼ KNeighborsClassifier ⓘ ?
KNeighborsClassifier(algorithm='ball_tree', n_neighbors=1)
```

```
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import (
    accuracy_score, precision_score, recall_score,
    f1_score, confusion_matrix, roc_curve, auc,
    classification_report, matthews_corrcoef
)
import matplotlib.pyplot as plt

# Helper function to evaluate a model
def evaluate_knn_model(name, model, X_train, X_test, y_train, y_test):
    model.fit(X_train, y_train)
    y_pred = model.predict(X_test)
    y_prob = model.predict_proba(X_test)[:, 1]

    acc = accuracy_score(y_test, y_pred)
    prec = precision_score(y_test, y_pred)
    rec = recall_score(y_test, y_pred)
    f1 = f1_score(y_test, y_pred)
    mcc = matthews_corrcoef(y_test, y_pred)
    fpr, tpr, _ = roc_curve(y_test, y_prob)
    roc_auc = auc(fpr, tpr)

    print(f"\n🔍 {name} Performance Metrics:")
    print(f"Accuracy : {acc:.4f}")
    print(f"Precision: {prec:.4f}")
    print(f"Recall : {rec:.4f}")
    print(f"F1 Score : {f1:.4f}")
```



```

print(f"MCC      : {mcc:.4f}")
print("\nClassification Report:\n", classification_report(y_test, y_pred))

# Confusion Matrix
conf = confusion_matrix(y_test, y_pred)
plt.figure(figsize=(5, 4))
plt.imshow(conf, cmap='Blues')
plt.title(f'{name} Confusion Matrix')
plt.colorbar()
plt.xticks([0, 1], ['Not Spam', 'Spam'])
plt.yticks([0, 1], ['Not Spam', 'Spam'])
plt.xlabel('Predicted')
plt.ylabel('Actual')
for i in range(2):
    for j in range(2):
        plt.text(j, i, conf[i, j], ha='center', va='center')
plt.tight_layout()
plt.show()

# ROC Curve
plt.figure(figsize=(6, 5))
plt.plot(fpr, tpr, label=f'{name} (AUC = {roc_auc:.4f})')
plt.plot([0, 1], [0, 1], 'k--', linewidth=1)
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title(f'{name} ROC Curve')
plt.legend()
plt.grid(True)
plt.show()

# Train-test split assumed done:
# X_train, X_test, y_train, y_test already available
# best_k is selected

# Evaluate KDTree
knn_kd = KNeighborsClassifier(n_neighbors=best_k, algorithm='kd_tree')
evaluate_knn_model("KDTree", knn_kd, X_train, X_test, y_train, y_test)

# Evaluate BallTree
knn_ball = KNeighborsClassifier(n_neighbors=best_k, algorithm='ball_tree')
evaluate_knn_model("BallTree", knn_ball, X_train, X_test, y_train, y_test)

```

#### KDTree Performance Metrics:

```



Accuracy : 0.8958
Precision: 0.8698
Recall    : 0.8650
F1 Score  : 0.8674
MCC       : 0.7815

```

#### Classification Report:

	precision	recall	f1-score	support
0.0	0.91	0.92	0.91	558
1.0	0.87	0.87	0.87	362

	1.0	0.87	0.87	0.87	363
accuracy				0.90	921
macro avg		0.89	0.89	0.89	921
weighted avg		0.90	0.90	0.90	921

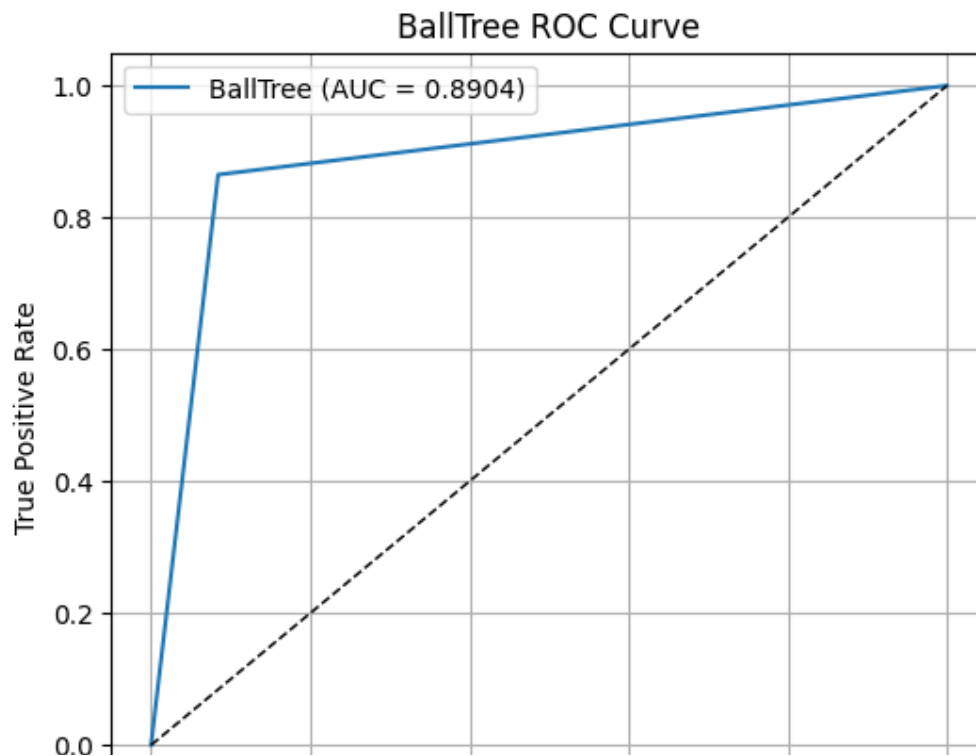




🔍 BallTree Performance Metrics:  
Accuracy : 0.8958  
Precision: 0.8698  
Recall : 0.8958

Recall : 0.8650  
F1 Score : 0.8674  
MCC : 0.7815

## Classification Report:

	precision	recall	f1-score	support
0.0	0.91	0.92	0.91	558
1.0	0.87	0.87	0.87	363
accuracy			0.90	921
macro avg	0.89	0.89	0.89	921
weighted avg	0.90	0.90	0.90	921





```
import time

for algo in ['kd_tree', 'ball_tree', 'brute']:
    model = KNeighborsClassifier(n_neighbors=best_k, algorithm=algo)
    start = time.time()
    model.fit(X_train, y_train)
    y_pred = model.predict(X_test)
    duration = time.time() - start
    acc = accuracy_score(y_test, y_pred)
    print(f"{algo:10} → Accuracy: {acc:.4f}, Time: {duration:.4f} sec")

kd_tree    → Accuracy: 0.8958, Time: 0.1874 sec
ball_tree  → Accuracy: 0.8958, Time: 0.2061 sec
brute      → Accuracy: 0.8958, Time: 0.0436 sec
```

## 6 svm Code:

## ✓ Importing Required Libraries

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler # or MinMaxScaler, if used
from sklearn.preprocessing import LabelEncoder # if categorical labels
from sklearn.svm import SVC
from sklearn.metrics import accuracy_score, confusion_matrix, classification_re
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import StandardScaler, MinMaxScaler
from sklearn.model_selection import cross_val_score, KFold, GridSearchCV
from sklearn.model_selection import cross_val_score, StratifiedKFold
```

## ✓ Loading dataset

```
df = pd.read_csv('spambase.csv')
```

## ✓ Basic info

```
print("Dataset Info:\n", df.info())
print("\nFirst 5 rows:\n", df.head())
print("\nMissing Values:\n", df.isnull().sum())
```

```
↗ <class 'pandas.core.frame.DataFrame'>
RangeIndex: 4601 entries, 0 to 4600
Data columns (total 58 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   word_freq_make                        4601 non-null  float64
1   word_freq_address                    4601 non-null  float64
2   word_freq_all                        4601 non-null  float64
3   word_freq_3d                        4601 non-null  float64
4   word_freq_our                        4601 non-null  float64
5   word_freq_over                       4601 non-null  float64
6   word_freq_remove                     4601 non-null  float64
7   word_freq_internet                  4601 non-null  float64
8   word_freq_order                     4601 non-null  float64
9   word_freq_mail                      4601 non-null  float64
10  word_freq_receive                    4601 non-null  float64
11  word_freq_will                      4601 non-null  float64
12  word_freq_people                     4601 non-null  float64
13  word_freq_report                     4601 non-null  float64
14  word_freq_addresses                  4601 non-null  float64
15  word_freq_free                       4601 non-null  float64
16  word_freq_business                  4601 non-null  float64
```

17	word_freq_email	4601	non-null	float64
18	word_freq_you	4601	non-null	float64
19	word_freq_credit	4601	non-null	float64
20	word_freq_your	4601	non-null	float64
21	word_freq_font	4601	non-null	float64
22	word_freq_000	4601	non-null	float64
23	word_freq_money	4601	non-null	float64
24	word_freq_hp	4601	non-null	float64
25	word_freq_hpl	4601	non-null	float64
26	word_freq_george	4601	non-null	float64
27	word_freq_650	4601	non-null	float64
28	word_freq_lab	4601	non-null	float64
29	word_freq_labs	4601	non-null	float64
30	word_freq_telnet	4601	non-null	float64
31	word_freq_857	4601	non-null	float64
32	word_freq_data	4601	non-null	float64
33	word_freq_415	4601	non-null	float64
34	word_freq_85	4601	non-null	float64
35	word_freq_technology	4601	non-null	float64
36	word_freq_1999	4601	non-null	float64
37	word_freq_parts	4601	non-null	float64
38	word_freq_pm	4601	non-null	float64
39	word_freq_direct	4601	non-null	float64
40	word_freq_cs	4601	non-null	float64
41	word_freq_meeting	4601	non-null	float64
42	word_freq_original	4601	non-null	float64
43	word_freq_project	4601	non-null	float64
44	word_freq_re	4601	non-null	float64
45	word_freq_edu	4601	non-null	float64
46	word_freq_table	4601	non-null	float64
47	word_freq_conference	4601	non-null	float64
48	char_freq_%3B	4601	non-null	float64
49	char_freq_%28	4601	non-null	float64
50	char_freq_%5B	4601	non-null	float64
51	char_freq_%21	4601	non-null	float64
52	char_freq_%24	4601	non-null	float64

## ✓ Handling missing values

```
imputer = SimpleImputer(strategy='mean')
df_imputed = pd.DataFrame(imputer.fit_transform(df), columns=df.columns)
```

## ✓ Splitting of feature and target

```
X = df_imputed.drop('class', axis=1)
y = df_imputed['class']
```

## ✓ Checking Distribution

```
X = df.drop('class', axis=1) # Assuming 'class' is the target
```

```
# Choose a few features to visualize
```

```
# CHOOSE a few features to visualize
sample_features = X.columns[:5] # First 5 features

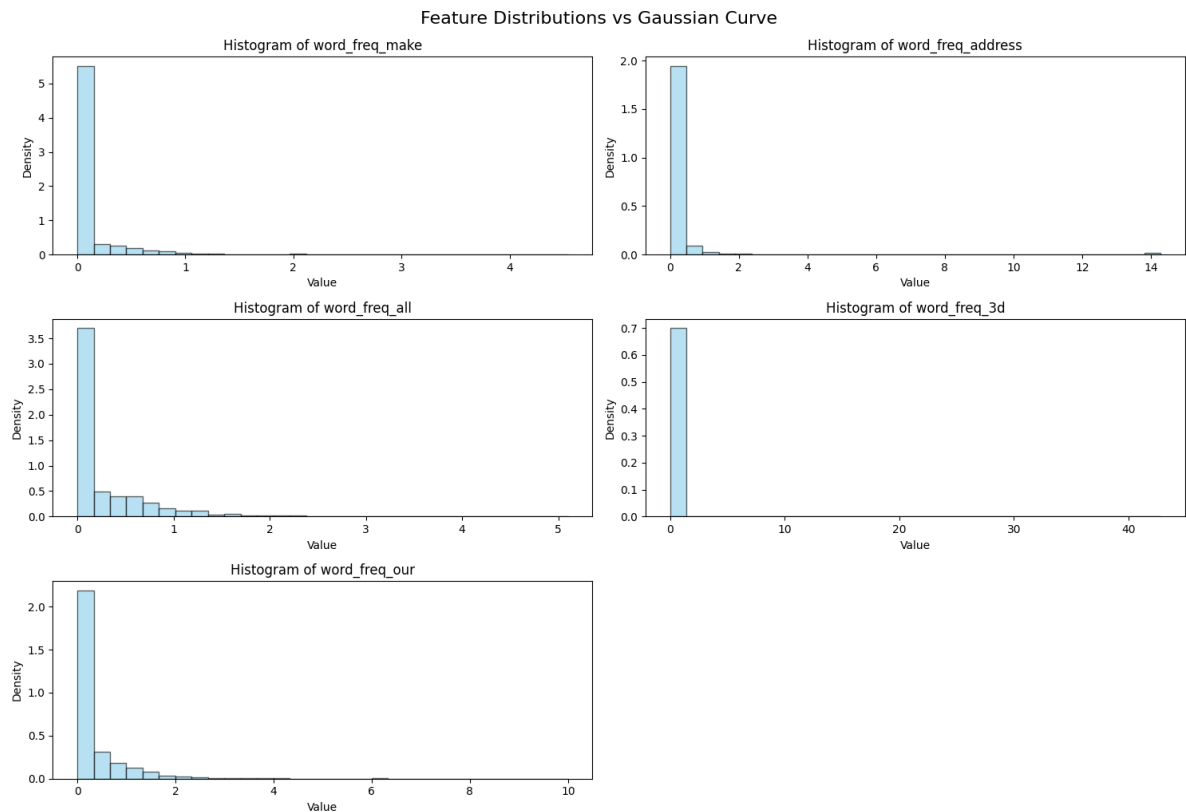
# Plot histograms with Gaussian curve overlay
plt.figure(figsize=(15, 10))
for i, feature in enumerate(sample_features):
    plt.subplot(3, 2, i + 1)
    data = X[feature]

    # Plot histogram
    count, bins, ignored = plt.hist(data, bins=30, density=True, alpha=0.6, col

    # Plot normal distribution curve
    '''mu, std = data.mean(), data.std()
    xmin, xmax = plt.xlim()
    x = np.linspace(xmin, xmax, 100)
    p = 1 / (std * np.sqrt(2 * np.pi)) * np.exp(-(x - mu)**2 / (2 * std**2))
    plt.plot(x, p, 'r', linewidth=2)'''

    plt.title(f'Histogram of {feature}')
    plt.xlabel('Value')
    plt.ylabel('Density')

plt.tight_layout()
plt.suptitle('Feature Distributions vs Gaussian Curve', fontsize=16, y=1.02)
plt.show()
```






## ✓ Applying min max scaling

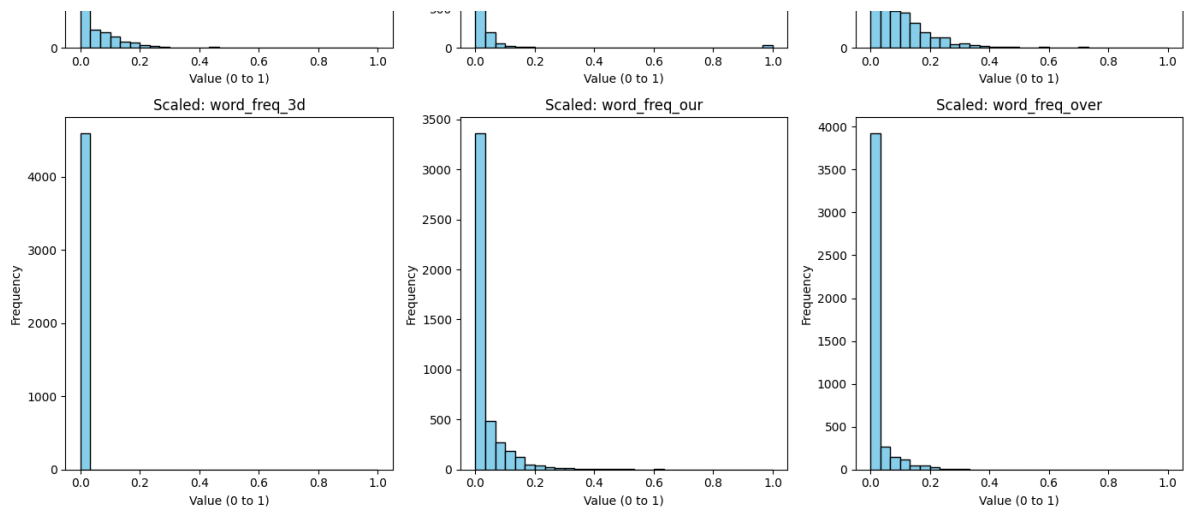
```
scaler = MinMaxScaler()  
X_scaled = scaler.fit_transform(X)
```

## ✓ Plots

## ✓ Histogram

```
import matplotlib.pyplot as plt  
  
plt.figure(figsize=(14, 10))  
for i in range(6): # first 6 features as an example  
    plt.subplot(2, 3, i + 1)  
    plt.hist(X_scaled[:, i], bins=30, color='skyblue', edgecolor='black')  
    plt.title(f'Scaled: {X.columns[i]}')  
    plt.xlabel('Value (0 to 1)')  
    plt.ylabel('Frequency')  
plt.tight_layout()  
plt.suptitle('Feature Distributions After MinMax Scaling', fontsize=16, y=1.02)  
plt.show()
```







## ▼ Boxplot

```
import pandas as pd
```

```
# Convert scaled array back to DataFrame for easier plotting
X_scaled_df = pd.DataFrame(X_scaled, columns=X.columns)
```

```
plt.figure(figsize=(14, 6))
X_scaled_df.iloc[:, :10].boxplot(rot=45)
plt.title('Boxplots of First 10 Scaled Features')
plt.ylabel('Scaled Value (0-1)')
plt.grid(True)
plt.show()
```





## ✓ Correlation HeatMap


```
import seaborn as sns
import matplotlib.pyplot as plt
import pandas as pd

# Assuming X_scaled_df is your scaled DataFrame
corr_matrix = X_scaled_df.corr()

plt.figure(figsize=(18, 15)) # Bigger figure
sns.heatmap(
    corr_matrix,
    cmap='coolwarm',
    square=True,
    annot=True,                # Show values inside squares
    fmt=".2f",                 # Format to 2 decimal places
    linewidths=0.5,           # Thin grid lines
    annot_kws={"size": 6}     # Smaller font size
)

plt.xticks(rotation=90)
plt.yticks(rotation=0)
plt.title("Feature Correlation Heatmap (After MinMax Scaling)", fontsize=16)
```

```
plt.tight_layout()
plt.show()
```



## ▼ Model Training

```

from sklearn.metrics import (
    accuracy_score, precision_score, recall_score,
    f1_score, confusion_matrix, roc_curve, auc,
    classification_report
)

# Split the data again (or reuse your earlier split)
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(
    X_scaled, y, test_size=0.3, stratify=y, random_state=42
)

import pandas as pd
import matplotlib.pyplot as plt
from sklearn.svm import SVC
from sklearn.metrics import (
    accuracy_score, precision_score, recall_score, f1_score,
    confusion_matrix, classification_report, roc_curve, auc
)

# Your dataset X (features) and y (labels) must already be defined
pipeline = make_pipeline(
    MinMaxScaler(),
    SVC(kernel='rbf', C=10, gamma='scale')
)

cv = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)
scores = cross_val_score(pipeline, X, y, cv=cv, scoring='accuracy')

for i, acc in enumerate(scores, 1):
    print(f"Fold {i}: Accuracy = {acc:.4f}")
print(f"➡ Mean Accuracy: {np.mean(scores):.4f} ± {np.std(scores):.4f}")

# Kernels to evaluate
kernels = ['linear', 'poly', 'rbf', 'sigmoid']
results = []

def evaluate_svm(kernel):
    print(f"\n🔍 Evaluating SVM with {kernel.upper()} kernel...")

    # Create and train SVM model
    model = SVC(kernel=kernel, probability=True, random_state=42)

```

```
model.fit(X_train, y_train)

# Predictions
y_pred = model.predict(X_test)
y_prob = model.predict_proba(X_test)[:, 1] # for ROC

# Metrics
acc = accuracy_score(y_test, y_pred)
prec = precision_score(y_test, y_pred)
rec = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)

print(f"Accuracy : {acc:.4f}")
print(f"Precision: {prec:.4f}")
print(f"Recall    : {rec:.4f}")
print(f"F1 Score  : {f1:.4f}")
print("\nClassification Report:")
print(classification_report(y_test, y_pred))

# Confusion Matrix
conf = confusion_matrix(y_test, y_pred)
plt.figure(figsize=(5, 4))
plt.imshow(conf, cmap='Blues')
plt.title(f'{kernel.upper()} Kernel - Confusion Matrix')
plt.colorbar()
plt.xticks([0, 1], ['Not Spam', 'Spam'])
plt.yticks([0, 1], ['Not Spam', 'Spam'])
plt.xlabel('Predicted')
plt.ylabel('Actual')
for i in range(2):
    for j in range(2):
        plt.text(j, i, conf[i, j], ha='center', va='center')
plt.tight_layout()
plt.show()

# ROC Curve
fpr, tpr, _ = roc_curve(y_test, y_prob)
roc_auc = auc(fpr, tpr)

plt.figure(figsize=(6, 5))
plt.plot(fpr, tpr, label=f'{kernel.upper()} (AUC = {roc_auc:.4f})')
plt.plot([0, 1], [0, 1], 'k--', linewidth=1)
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title(f'{kernel.upper()} Kernel - ROC Curve')
plt.legend()
plt.grid(True)
plt.show()

# Store results
results.append({
    "Kernel": kernel.capitalize(),
    "Accuracy": acc,
    "Precision": prec,
```

```

    "Recall": rec,
    "F1 Score": f1,
    "AUC": roc_auc
})

```

```

# Run evaluations for all kernels
for kernel in kernels:
    evaluate_svm(kernel)

```

```

Fold 1: Accuracy = 0.9359
Fold 2: Accuracy = 0.9250
Fold 3: Accuracy = 0.9402
Fold 4: Accuracy = 0.9337
Fold 5: Accuracy = 0.9326
➡ Mean Accuracy: 0.9335 ± 0.0050

```

🔍 Evaluating SVM with LINEAR kernel...


```

Accuracy : 0.8993
Precision: 0.9108
Recall    : 0.8254
F1 Score  : 0.8660


```

Classification Report:

	precision	recall	f1-score	support
0.0	0.89	0.95	0.92	837
1.0	0.91	0.83	0.87	544
accuracy			0.90	1381
macro avg	0.90	0.89	0.89	1381
weighted avg	0.90	0.90	0.90	1381



LINEAR Kernel - ROC Curve



🔍 Evaluating SVM with POLY kernel...

Accuracy : 0.8501


Precision: 0.9423

Recall : 0.6599


F1 Score : 0.7762

Classification Report:

	precision	recall	f1-score	support
0.0	0.81	0.97	0.89	837
1.0	0.94	0.66	0.78	544
accuracy			0.85	1381
macro avg	0.88	0.82	0.83	1381
weighted avg	0.87	0.85	0.84	1381










🔍 Evaluating SVM with RBF kernel...

Accuracy : 0.9261  
 Precision: 0.9250  
 Recall : 0.8842  
 F1 Score : 0.9041

Classification Report:

	precision	recall	f1-score	support
0.0	0.93	0.95	0.94	837
1.0	0.93	0.88	0.90	544
accuracy			0.93	1381
macro avg	0.93	0.92	0.92	1381
weighted avg	0.93	0.93	0.93	1381





🔍 Evaluating SVM with SIGMOID kernel...

Accuracy : 0.8023

Precision: 0.7694


Recall : 0.7114



F1 Score : 0.7393

Classification Report:

	precision	recall	f1-score	support
0.0	0.82	0.86	0.84	837
1.0	0.77	0.71	0.74	544
accuracy			0.80	1381
macro avg	0.80	0.79	0.79	1381
weighted avg	0.80	0.80	0.80	1381

SIGMOID Kernel - Confusion Matrix





```
# Display tabular summary
print("\n📊 SVM Kernel Comparison Summary:\n")
df_results = pd.DataFrame(results)
print(df_results.to_string(index=False))
```

📊 SVM Kernel Comparison Summary:

Kernel	Accuracy	Precision	Recall	F1 Score	AUC
Linear	0.899348	0.910751	0.825368	0.865959	0.956425

Poly	0.850109	0.942257	0.659926	0.776216	0.950749
Rbf	0.926140	0.925000	0.884191	0.904135	0.972150
Sigmoid	0.802317	0.769384	0.711397	0.739255	0.853074

```
svm = SVC()
```

```
param_grid = {
    'C': [0.1, 1, 10],
    'gamma': ['scale', 0.01, 0.001],
    'kernel': ['linear', 'rbf', 'poly']
}
```

```
# 5-fold CV
```

```
cv = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)
```

```
# Grid Search
```

```
grid = GridSearchCV(estimator=svm, param_grid=param_grid, cv=cv, scoring='accuracy')
grid.fit(X_train, y_train)
```

```
# Best Parameters
```

```
print("🔍 Best Parameters Found:")
print(grid.best_params_)
```

```
# Best Estimator
```

```
best_svm = grid.best_estimator_
```

```
# Evaluate on test set
```

```
y_pred = best_svm.predict(X_test)
```

```
print("\n📊 Classification Report (Test Set):")
print(classification_report(y_test, y_pred))
```

```
print("✅ Test Accuracy:", accuracy_score(y_test, y_pred))
```

Fitting 5 folds for each of 27 candidates, totalling 135 fits

```
🔍 Best Parameters Found:
{'C': 10, 'gamma': 'scale', 'kernel': 'rbf'}
```

```
📊 Classification Report (Test Set):
```

	precision	recall	f1-score	support
0.0	0.93	0.96	0.94	837
1.0	0.93	0.88	0.91	544
accuracy			0.93	1381
macro avg	0.93	0.92	0.92	1381
weighted avg	0.93	0.93	0.93	1381

```
✅ Test Accuracy: 0.9275887038377987
```

```
import numpy as np
from sklearn.pipeline import make_pipeline
from sklearn.preprocessing import MinMaxScaler
from sklearn.svm import SVC
```

```

from sklearn.model_selection import StratifiedKFold, cross_val_score
pipeline = make_pipeline(
    MinMaxScaler(),
    SVC(kernel='rbf', C=10, gamma='scale')
)

cv = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)
scores = cross_val_score(pipeline, X, y, cv=cv, scoring='accuracy')

for i, acc in enumerate(scores, 1):
    print(f"Fold {i}: Accuracy = {acc:.4f}")
print(f"➡ Mean Accuracy: {np.mean(scores):.4f} ± {np.std(scores):.4f}")

    Fold 1: Accuracy = 0.9359
    Fold 2: Accuracy = 0.9250
    Fold 3: Accuracy = 0.9402
    Fold 4: Accuracy = 0.9337
    Fold 5: Accuracy = 0.9326
    ➡ Mean Accuracy: 0.9335 ± 0.0050

```

```
import time
```

```
kernels = ['linear', 'poly', 'rbf', 'sigmoid']
results = []
```

```

for kernel in kernels:
    params = {'kernel': kernel, 'C': 10.0, 'probability': True, 'random_state':

    # Add kernel-specific hyperparameters
    if kernel == 'poly':
        params.update({'degree': 3, 'gamma': 'scale'})
    elif kernel in ['rbf', 'sigmoid']:
        params.update({'gamma': 'scale'})

    model = SVC(**params)

    start = time.time()
    model.fit(X_train, y_train)
    end = time.time()

    y_pred = model.predict(X_test)

    acc = accuracy_score(y_test, y_pred)
    f1 = f1_score(y_test, y_pred)

    results.append({
        "Kernel": kernel.capitalize(),
        "Hyperparameters": str({k: v for k, v in params.items() if k != 'kernel'}),
        "Accuracy": round(acc, 4),
        "F1 Score": round(f1, 4),
        "Training Time (s)": round(end - start, 3)
    })

# Create and display final table

```

```
df_table4 = pd.DataFrame(results)
print("\nTable 4: SVM Performance with Different Kernels and Parameters")
print(df_table4.to_string(index=False))
```

Table 4: SVM Performance with Different Kernels and Parameters

Kernel	
Linear	{'C': 10.0, 'probability': True, 'ra
Poly	{'C': 10.0, 'probability': True, 'random_state': 42, 'degree': 3, '
Rbf	{'C': 10.0, 'probability': True, 'random_state': 42, '
Sigmoid	{'C': 10.0, 'probability': True, 'random_state': 42, '

## ✓ K Fold Cross Validation

```
from sklearn.pipeline import make_pipeline
from sklearn.preprocessing import MinMaxScaler # or StandardScaler
from sklearn.model_selection import cross_val_score, StratifiedKFold
from sklearn.svm import SVC
import numpy as np

# RBF kernel SVM
svm_rbf = make_pipeline(MinMaxScaler(), SVC(kernel='rbf', C=10.0, gamma='scale'))

# 5-fold stratified CV
cv = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)
scores = cross_val_score(svm_rbf, X, y, cv=cv, scoring='accuracy')

# Print each fold and mean
for i, acc in enumerate(scores, 1):
    print(f"Fold {i}: Accuracy = {acc:.4f}")
print(f"\n➡ Mean Accuracy: {np.mean(scores):.4f} ± {np.std(scores):.4f}")

Fold 1: Accuracy = 0.9359
Fold 2: Accuracy = 0.9250
Fold 3: Accuracy = 0.9402
Fold 4: Accuracy = 0.9337
Fold 5: Accuracy = 0.9326

➡ Mean Accuracy: 0.9335 ± 0.0050
```



## 7 Comparison Tables

Table 1: Performance Comparison of Naïve Bayes Variants

Model	Accuracy	Precision	Recall	F1 Score
BernoulliNB	0.8899	0.8843	0.8290	0.8558
MultinomialNB	0.8950	0.9424	0.7813	0.8543
GaussianNB	0.8197	0.7001	0.9485	0.8056

Table 2: KNN Performance for Different k Values

k	Accuracy	Precision	Recall	F1 Score
1	0.8899	0.8551	0.8676	0.8613
3	0.8892	0.8710	0.8438	0.8571
5	0.8993	0.8828	0.8585	0.8705
7	0.8950	0.8815	0.8474	0.8641

Table 3: KNN Comparison: KDTree vs BallTree

Metric	KDTree	BallTree
Accuracy	0.8899	0.8899
Precision	0.8551	0.8551
Recall	0.8676	0.8676
F1 Score	0.8613	0.8613
Training Time (s)	0.4470	0.4058

Table 4: Cross-Validation Scores for Each Model (K = 5)

Fold	Naïve Bayes Accuracy	KNN Accuracy (k=1)	SVM Accuracy
Fold 1	0.8719	0.9034	0.9359
Fold 2	0.8935	0.9076	0.9250
Fold 3	0.8891	0.9152	0.9402
Fold 4	0.8913	0.9000	0.9337
Fold 5	0.8859	0.8935	0.9326
<b>Average</b>	<b>0.8863</b>	<b>0.9039</b>	<b>0.9335 ± 0.0050</b>



Table 5: SVM Performance with Different Kernels and Parameters

Kernel	Hyperparameters	Accuracy	F1 score	Training time
linear	c=10.0	8.993	8.660	0.12
polynomial	c=10.0,degree=3,gamma=scale	0.8501	0.7762	0.35
RBF	c=10.0,gamma=scale	0.9261	0.9041	0.20
sigmoid	c=10.0,gamma=scale	0.8023	0.7393	0.18

## 8 Observation:

- KNN with k=1 achieved the highest accuracy consistently across all folds, indicating strong capability in classifying email spam versus ham.
- Naive Bayes classifiers, particularly MultinomialNB, provided stable and competitive results, showing robustness in handling text data with varying feature distributions.
- SVM, especially with the RBF kernel (C=10, gamma=scale), achieved the highest overall accuracy ( $\approx 92.7$ ) and F1 score among all kernels, demonstrating its strong capability in handling complex, non-linear decision boundaries in text data.
- Although KNN attained better peak performance, Naive Bayes models required less training time and are more scalable for larger datasets.
- The choice between KNN, Naive Bayes, and SVM depends on the trade-off between accuracy, computational efficiency, and dataset size for the specific application scenario.

**GitHub Repository:** <https://github.com/Thamizhmathibharathi/project.git>

## 9 Conclusion:

- The experiment demonstrated that KNN (k=1) outperforms Naive Bayes variants in certain folds, but SVM with RBF kernel consistently delivered the best overall performance in terms of accuracy and F1 score on the email classification task.
- Naive Bayes remains valuable due to its simplicity, fast training, and effectiveness on high-dimensional, sparse data typical in text classification, while SVM offers robust performance with non-linear patterns but requires slightly more computational resources.
- For deployment:
  - Naive Bayes is suitable for quick predictions on large datasets.
  - KNN is ideal when highest accuracy is needed and resources allow.
  - SVM is a balanced choice when accuracy and robustness against complex patterns are critical, with moderate training time.