

Sri Sivasubramaniya Nadar College of Engineering, Chennai
(An autonomous Institution affiliated to Anna University)

Degree & Branch	B.E. Computer Science & Engineering	Semester	V
Subject Code & Name	ICS1512 & Machine Learning Algorithms Laboratory		
Academic year	2025-2026 (Odd)	Batch:2023-2028	Due date:

Experiment 2 : Loan Amount Prediction using Linear Regression

1 Aim:

To develop and evaluate a machine learning model using Python to predict loan amounts based on applicant financial and credit-related features, utilizing Linear Regression and K-Fold Cross-Validation for performance assessment.

2 Libraries used:

- Numpy
- Pandas
- Matplotlib
- Scikit-learn
- Seaborn

3 Objective:

- To preprocess the loan dataset by handling missing values, encoding categorical variables, selecting relevant features, and splitting the data for training, validation, and testing purposes.
- To build and evaluate a Linear Regression model using K-Fold Cross-Validation, analyze performance metrics such as MSE, RMSE, MAE, and R^2 , and interpret the results through residual and prediction plots to assess model effectiveness.

4 Mathematical/theoretical description of the algorithm/objective performed:

4.1 Handling Missing Values

Missing values can negatively impact the performance of machine learning models by:

- Distorting statistical summaries
- Causing errors in training algorithms
- Leading to biased predictions

So, it's crucial to detect and properly handle them before modeling.

There are several ways to handle missing values:

- Missing values in a dataset can be handled using imputation techniques such as replacing them with the mean, median, or mode of the respective feature by using `fillna()` method in pandas
- If a column contains a large number of missing values and does not contribute significantly to the prediction task, it can be dropped to simplify the dataset. Removing such irrelevant or incomplete features helps reduce noise and improve model efficiency.

4.2 Label encoding:

To train machine learning models, all input features must be in numeric format. Hence, categorical variables (like "Yes"/"No", "Graduate"/"Not Graduate") need to be converted into numbers. This process is essential for enabling algorithms to interpret and process the data correctly.

- Categorical values can be directly replaced with numeric codes, such as mapping "Yes" to 1 and "No" to 0. This is useful when the categories are binary or have no specific order. It ensures compatibility with machine learning models that require numerical input.
- If a categorical feature has more than two values, simple replacement may introduce unintended ordinal relationships. In such cases, **one-hot encoding** is preferred, where each category becomes a separate binary column. This prevents the model from assuming any order or ranking between the categories.

4.3 Plotting:

To better understand the patterns and relationships within the dataset, various data visualization techniques are used. These plots help in identifying correlations, distributions, and outliers effectively.

- The **heatmap()** function visualizes the correlation between numerical features using a colored matrix. Darker shades typically represent stronger correlations, helping identify redundant or related features. It's a useful tool for understanding feature interdependencies before model building.

- A **histogram** displays the frequency distribution of a numeric variable, showing how data is spread across ranges. It helps detect skewness, modality, and presence of outliers or missing value gaps. This is often used as a first step in understanding individual feature behavior.
- A **box plot** (or whisker plot) shows the spread and central tendency of a feature using quartiles. It clearly identifies the median, interquartile range (IQR), and outliers. This makes it an excellent tool for spotting extreme values and data symmetry.

4.4 Standardization:

- Standardization is a feature scaling technique that **transforms data to have a mean of 0 and a standard deviation of 1**. It is especially useful when features have different units or scales, ensuring all variables contribute equally to the model. Many machine learning algorithms, like logistic regression and KNN, perform better when data is standardized.
- The formula used for standardization is:

$$z = \frac{x - \mu}{\sigma}$$

- where x is the original value, μ is the mean, and σ is the standard deviation. This process centers the data around zero and makes it easier for models to converge efficiently.

The preprocessing steps involved handling missing values, encoding categorical variables, visualizing data using heatmaps, histograms, and boxplots, and addressing outliers. Additionally, feature standardization was applied to bring all variables to a common scale, ensuring better model performance and stability.

5 Code:

loan_amount_prediction

July 29, 2025

```
[2]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
[3]: df=pd.read_csv('train.csv')
df.head()
```

```
[3]: Customer ID      Name Gender  Age  Income (USD)  Income Stability \
0      C-36995  Frederica Shealy    F   56      1933.05             Low
1      C-33999  America Calderone    M   32      4952.91             Low
2      C-3770   Rosetta Verne      F   65       988.19             High
3      C-26480    Zoe Chitty      F   65         NaN             High
4      C-23459   Afton Venema      F   31      2614.77             Low
```

```
      Profession      Type of Employment      Location  Loan Amount Request (USD) \
0      Working      Sales staff  Semi-Urban      72809.58
1      Working      NaN  Semi-Urban      46837.47
2  Pensioner      NaN  Semi-Urban      45593.04
3  Pensioner      NaN      Rural      80057.92
4      Working  High skill tech staff  Semi-Urban      113858.89
```

```
      ...  Credit Score  No. of Defaults  Has Active Credit Card  Property ID \
0      ...      809.44      0      NaN      746
1      ...      780.40      0      Unpossessed      608
2      ...      833.15      0      Unpossessed      546
3      ...      832.70      1      Unpossessed      890
4      ...      745.55      1      Active      715
```

```
      Property Age  Property Type  Property Location  Co-Applicant \
0      1933.05      4      Rural      1
1      4952.91      2      Rural      1
2      988.19      2      Urban      0
3      NaN      2      Semi-Urban      1
4      2614.77      4      Semi-Urban      1
```

```
      Property Price  Loan Sanction Amount (USD)
0      119933.46      54607.18
```

1	54791.00	37469.98
2	72440.58	36474.43
3	121441.51	56040.54
4	208567.91	74008.28

[5 rows x 24 columns]

```
[4]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 30000 entries, 0 to 29999
Data columns (total 24 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Customer ID                          30000 non-null  object
1   Name                                30000 non-null  object
2   Gender                              29947 non-null  object
3   Age                                 30000 non-null  int64
4   Income (USD)                        25424 non-null  float64
5   Income Stability                     28317 non-null  object
6   Profession                           30000 non-null  object
7   Type of Employment                  22730 non-null  object
8   Location                            30000 non-null  object
9   Loan Amount Request (USD)           30000 non-null  float64
10  Current Loan Expenses (USD)         29828 non-null  float64
11  Expense Type 1                      30000 non-null  object
12  Expense Type 2                      30000 non-null  object
13  Dependents                          27507 non-null  float64
14  Credit Score                        28297 non-null  float64
15  No. of Defaults                     30000 non-null  int64
16  Has Active Credit Card              28434 non-null  object
17  Property ID                        30000 non-null  int64
18  Property Age                       25150 non-null  float64
19  Property Type                      30000 non-null  int64
20  Property Location                   29644 non-null  object
21  Co-Applicant                       30000 non-null  int64
22  Property Price                     30000 non-null  float64
23  Loan Sanction Amount (USD)         29660 non-null  float64
dtypes: float64(8), int64(5), object(11)
memory usage: 5.5+ MB
```

```
[5]: df['Co-Applicant'].unique()
```

```
[5]: array([ 1, 0, -999])
```

```
[6]: df['Has Active Credit Card'].unique()
```

```
[6]: array([nan, 'Unpossessed', 'Active', 'Inactive'], dtype=object)
```

```
[7]: #Removing unnecessary columns
df = df.drop(columns=["Customer ID", "Name"])
```

```
[8]: # Replace -999 with NaN
df['Co-Applicant'] = df['Co-Applicant'].replace(-999, np.nan)

# Option 1: Impute missing values (e.g., assume no co-applicant)
df['Co-Applicant'] = df['Co-Applicant'].fillna(0)
```

```
[9]: # Fill NaN with 'Unknown'
df['Has Active Credit Card'] = df['Has Active Credit Card'].fillna('Unknown')

# Optional: Encode as ordinal
credit_card_map = {
    'Unpossessed': 0,
    'Inactive': 1,
    'Active': 2,
    'Unknown': -1
}
df['Has Active Credit Card'] = df['Has Active Credit Card'].map(credit_card_map)
```

```
[10]: df.isnull().sum()
```

```
[10]: Gender                53
Age                        0
Income (USD)              4576
Income Stability          1683
Profession                0
Type of Employment       7270
Location                  0
Loan Amount Request (USD) 0
Current Loan Expenses (USD) 172
Expense Type 1            0
Expense Type 2            0
Dependents                2493
Credit Score              1703
No. of Defaults           0
Has Active Credit Card     0
Property ID               0
Property Age              4850
Property Type             0
Property Location         356
Co-Applicant              0
Property Price            0
Loan Sanction Amount (USD) 340
dtype: int64
```

```
[11]: #Filling null values
df['Gender']=df['Gender'].fillna(df['Gender'].mode()[0])
df['Income (USD)']=df['Income (USD)'].fillna(df['Income (USD)'].median())
df['Income Stability']=df['Income Stability'].fillna(df['Income Stability'].
↳mode()[0])

[12]: #Dropping this column due to presence of more null values and may categories
df['Type of Employment'].unique()
df=df.drop(columns=['Type of Employment'])

[13]: #Current Loan Expenses (USD) - Numeric → fill with median
df['Current Loan Expenses (USD)'] = df['Current Loan Expenses (USD)'].
↳fillna(df['Current Loan Expenses (USD)'].median())

#Dependents - Numeric → fill with mode (likely a small integer like 1 or 2)
df['Dependents'] = df['Dependents'].fillna(df['Dependents'].mode()[0])

#Credit Score - Numeric → fill with median
df['Credit Score'] = df['Credit Score'].fillna(df['Credit Score'].median())

#Property Age - Numeric → fill with median
df['Property Age'] = df['Property Age'].fillna(df['Property Age'].median())

#Property Location - Categorical → fill with mode
df['Property Location'] = df['Property Location'].fillna(df['Property_
↳Location'].mode()[0])

# Loan Sanction Amount (USD) - Numeric → fill with median
df['Loan Sanction Amount (USD)'] = df['Loan Sanction Amount (USD)'].
↳fillna(df['Loan Sanction Amount (USD)'].median())

df['Loan Sanction Amount (USD)'] = df['Loan Sanction Amount (USD)'].replace(0,↳
↳df['Loan Sanction Amount (USD)'].median())

[14]: df.isnull().sum()
```

```
[14]: Gender          0
Age                  0
Income (USD)         0
Income Stability     0
Profession           0
Location             0
Loan Amount Request (USD) 0
Current Loan Expenses (USD) 0
Expense Type 1       0
Expense Type 2       0
Dependents           0
```

Credit Score	0
No. of Defaults	0
Has Active Credit Card	0
Property ID	0
Property Age	0
Property Type	0
Property Location	0
Co-Applicant	0
Property Price	0
Loan Sanction Amount (USD)	0
dtype: int64	

Encoding of variables with values

```
[15]: from sklearn.preprocessing import LabelEncoder

# List of categorical columns
cat_cols = [
    'Gender', 'Income Stability', 'Profession',
    'Expense Type 1', 'Expense Type 2',
    'Has Active Credit Card', 'Property Type', 'Property Location', 'Location'
]

# Create a label encoder instance
le = LabelEncoder()

# Apply label encoding to each column
for col in cat_cols:
    df[col] = le.fit_transform(df[col])
```

Standardization of Features

```
[16]: from sklearn.preprocessing import StandardScaler

# Identify numeric columns (excluding categorical and target)
numeric_cols = df.select_dtypes(include=['float64', 'int64']).columns.tolist()

# Optionally exclude target column (e.g., 'Loan Sanction Amount (USD)')
numeric_cols.remove('Loan Sanction Amount (USD)')

# Initialize scaler
scaler = StandardScaler()

# Fit and transform numeric features
df[numeric_cols] = scaler.fit_transform(df[numeric_cols])
```

```
[17]: df.head(10)
```


[17]:

	Gender	Age	Income (USD)	Income Stability	Profession	Location	\
0	-1.007092	0.991451	-0.061266	0.305833	0.834973	0.142149	
1	0.992958	-0.504355	0.229972	0.305833	0.834973	0.142149	
2	-1.007092	1.552379	-0.152389	-3.269763	-0.686548	0.142149	
3	-1.007092	1.552379	-0.033357	-3.269763	-0.686548	-1.762481	
4	-1.007092	-0.566680	0.004480	0.305833	0.834973	0.142149	
5	-1.007092	1.240752	-0.128594	0.305833	-0.306168	-1.762481	
6	0.992958	0.181223	-0.019940	0.305833	0.834973	0.142149	
7	-1.007092	0.305874	-0.033357	0.305833	-0.306168	0.142149	
8	-1.007092	-0.130403	-0.122697	0.305833	0.834973	-1.762481	
9	0.992958	-1.376908	-0.098577	0.305833	0.834973	-1.762481	

	Loan Amount Request (USD)	Current Loan Expenses (USD)	Expense Type 1	\
0	-0.269027	-0.660358	-0.749241	
1	-0.705269	0.392886	-0.749241	
2	-0.726171	-0.946193	-0.749241	
3	-0.147279	-0.422775	-0.749241	
4	0.420461	0.374693	-0.749241	
5	-0.913593	-0.906788	-0.749241	
6	1.070530	1.227526	1.334685	
7	2.544436	1.682224	-0.749241	
8	-0.901713	-1.012348	-0.749241	
9	-0.784989	0.411038	-0.749241	

	Expense Type 2	...	Credit Score	No. of Defaults	Has Active Credit Card	\
0	-1.433524	...	0.992493	-0.490502	-2.096903	
1	0.697582	...	0.578136	-0.490502	-1.001762	
2	0.697582	...	1.330799	-0.490502	-1.001762	
3	0.697582	...	1.324379	2.038728	-1.001762	
4	0.697582	...	0.080879	2.038728	1.188520	
5	-1.433524	...	-0.795636	2.038728	0.093379	
6	0.697582	...	-1.463830	-0.490502	-1.001762	
7	-1.433524	...	1.032730	-0.490502	1.188520	
8	0.697582	...	-0.493571	2.038728	1.188520	
9	-1.433524	...	-1.806987	-0.490502	-1.001762	

	Property ID	Property Age	Property Type	Property Location	Co-Applicant	\
0	0.846998	-0.060969	1.376731	-1.214540	0.419205	
1	0.368086	0.230298	-0.411309	-1.214540	0.419205	
2	0.152923	-0.152102	-0.411309	1.283229	-2.385467	
3	1.346732	-0.032979	-0.411309	0.034344	0.419205	
4	0.739417	0.004783	1.376731	0.034344	0.419205	
5	-0.037948	-0.128305	-0.411309	-1.214540	0.419205	
6	-0.954127	-0.019639	-1.305329	0.034344	0.419205	
7	-0.652204	-0.032979	-0.411309	1.283229	0.419205	
8	-0.905541	-0.122407	1.376731	-1.214540	0.419205	
9	1.322440	-0.098284	-0.411309	1.283229	0.419205	

	Property Price	Loan Sanction Amount (USD)
0	-0.126419	54607.180
1	-0.822772	37469.980
2	-0.634103	36474.430
3	-0.110298	56040.540
4	0.821057	74008.280
5	-0.947245	22382.570
6	0.954495	35209.395
7	2.878533	168218.240
8	-0.821570	22842.290
9	-0.681642	35209.395

[10 rows x 21 columns]



EDA



Histogram

```
[18]: import seaborn as sns
import matplotlib.pyplot as plt

# Plot distributions for selected numeric columns
cols_to_plot = ['Income (USD)', 'Credit Score', 'Loan Amount Request (USD)',
               ↪ 'Loan Sanction Amount (USD)']

for col in cols_to_plot:
    plt.figure(figsize=(6, 4))
    sns.histplot(df[col], kde=True, bins=30)
    plt.title(f'Distribution of {col}')
    plt.tight_layout()
    plt.show()
```

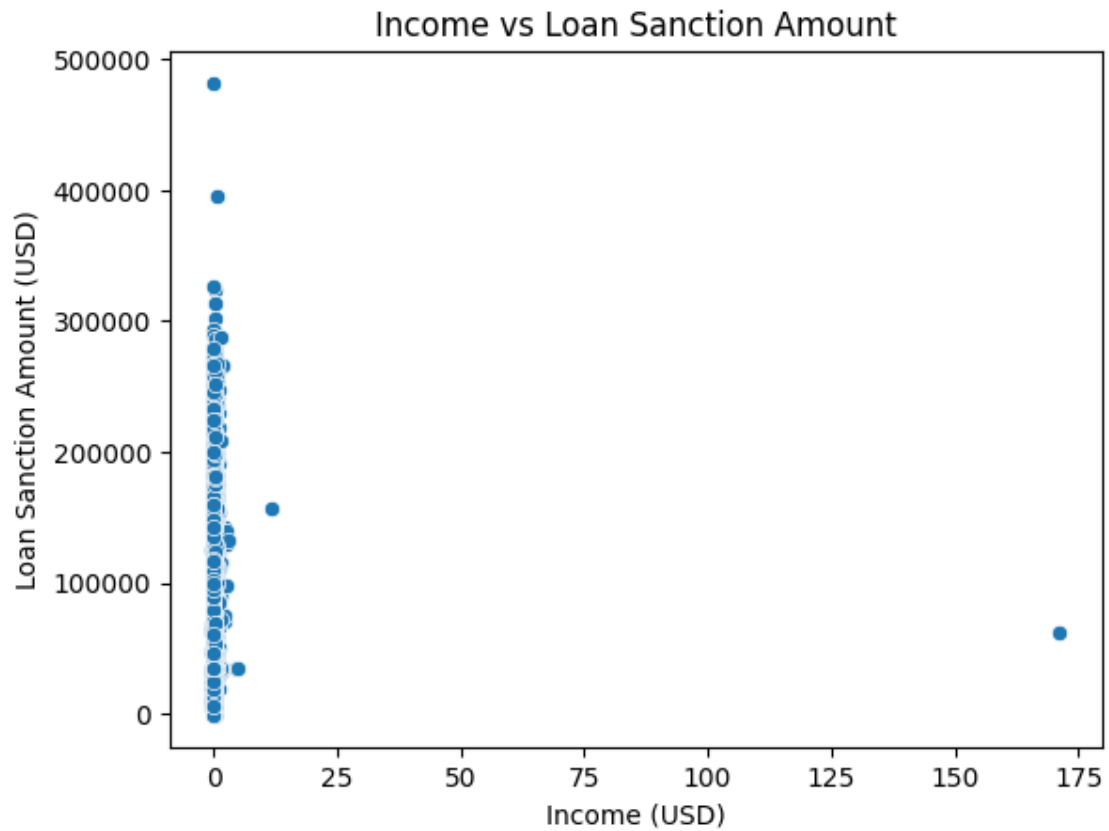





Scatter Plot

```
[19]: # Income vs Loan Sanction Amount
sns.scatterplot(data=df, x='Income (USD)', y='Loan Sanction Amount (USD)')
plt.title('Income vs Loan Sanction Amount')
plt.show()


# Credit Score vs Loan Sanction Amount
sns.scatterplot(data=df, x='Credit Score', y='Loan Sanction Amount (USD)')
plt.title('Credit Score vs Loan Sanction Amount')
plt.show()
```





Correlation heatmap

```
[20]: plt.figure(figsize=(12, 8))
sns.heatmap(df.corr(), annot=True, cmap='coolwarm', fmt=".2f")
plt.title('Correlation Heatmap')
plt.show()
```





BoxPlot

```
[21]: # Boxplot for Income
sns.boxplot(x=df['Income (USD)'])
plt.title('Boxplot of Income')
plt.show()

# Boxplot for Loan Amount Request
sns.boxplot(x=df['Loan Amount Request (USD)'])
plt.title('Boxplot of Loan Amount Request')
plt.show()
```

Boxplot of Income





Train Test Split

```
[22]: from sklearn.model_selection import train_test_split

# Define target variable
target = 'Loan Sanction Amount (USD)'

# Define feature columns
X = df.drop(columns=[target])
y = df[target]

# Split into train and test sets (80% train, 20% test)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    ↪random_state=42)
```

Model Training

```
[23]: from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score, mean_absolute_error

# Initialize and train the model
```

```

model = LinearRegression()
model.fit(X_train, y_train)

# Predict on test set
y_pred = model.predict(X_test)

# Evaluation Metrics
mse = mean_squared_error(y_test, y_pred)
rmse = np.sqrt(mse)
mae = mean_absolute_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

# Print results
print(f"Mean Squared Error (MSE): {mse:.2f}")
print(f"Root Mean Squared Error (RMSE): {rmse:.2f}")
print(f"Mean Absolute Error (MAE): {mae:.2f}")
print(f"R2 Score: {r2:.2f}")

```

Mean Squared Error (MSE): 527445271.77
 Root Mean Squared Error (RMSE): 22966.18
 Mean Absolute Error (MAE): 13803.42
 R² Score: 0.69

```

[24]: from sklearn.model_selection import KFold, cross_val_score
import numpy as np

# Define K-Fold with 5 splits
kf = KFold(n_splits=5, shuffle=True, random_state=42)

# Custom scoring functions
mse_scores = cross_val_score(model, X, y, scoring='neg_mean_squared_error',
                               cv=kf)
mae_scores = cross_val_score(model, X, y, scoring='neg_mean_absolute_error',
                               cv=kf)
r2_scores = cross_val_score(model, X, y, scoring='r2', cv=kf)

# Convert negative MSE/MAE to positive
mse_scores = -mse_scores
mae_scores = -mae_scores
rmse_scores = np.sqrt(mse_scores)

# Print metrics per fold
print("Fold-wise Metrics:")
for i in range(len(mse_scores)):
    print(f"Fold {i+1}:")
    print(f"    MSE : {mse_scores[i]:.2f}")

```

```

print(f" RMSE: {rmse_scores[i]:.2f}")
print(f" MAE : {mae_scores[i]:.2f}")
print(f" R2 : {r2_scores[i]:.2f}")
print()

# Print average performance
print("Average Metrics Across Folds:")
print(f"Average MSE : {mse_scores.mean():.2f}")
print(f"Average RMSE: {rmse_scores.mean():.2f}")
print(f"Average MAE : {mae_scores.mean():.2f}")
print(f"Average R2 : {r2_scores.mean():.2f}")

```

Fold-wise Metrics:

Fold 1:

```

MSE : 527445271.77
RMSE: 22966.18
MAE : 13803.42
R2 : 0.69

```

Fold 2:

```

MSE : 493351608.13
RMSE: 22211.52
MAE : 13779.90
R2 : 0.70

```

Fold 3:

```

MSE : 544801753.47
RMSE: 23340.99
MAE : 14030.71
R2 : 0.67

```

Fold 4:

```

MSE : 513654615.80
RMSE: 22663.95
MAE : 14044.93
R2 : 0.70

```

Fold 5:

```

MSE : 440761214.18
RMSE: 20994.31
MAE : 13347.16
R2 : 0.73

```

Average Metrics Across Folds:

```

Average MSE : 504002892.67
Average RMSE: 22435.39
Average MAE : 13801.23
Average R2 : 0.70


```

Actual vs Predicted values

```
[25]: import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.metrics import r2_score

# Predict values
y_pred = model.predict(X_test)

# Plot
plt.figure(figsize=(8, 6))
sns.scatterplot(x=y_test, y=y_pred, alpha=0.7)
plt.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()],
         color='red', linestyle='--')
plt.xlabel("Actual Loan Sanction Amount")
plt.ylabel("Predicted Loan Sanction Amount")
plt.title(f"Actual vs Predicted ( $R^2 = {r2\_score(y\_test, y\_pred):.2f}$ )")
plt.grid(True)
plt.show()
```




```
[26]: # Create a DataFrame for coefficients
coeff_df = pd.DataFrame({
    'Feature': X_train.columns,
    'Coefficient': model.coef_
}).sort_values(by='Coefficient', key=abs, ascending=False)

# Plot
plt.figure(figsize=(10, 6))
sns.barplot(x='Coefficient', y='Feature', data=coeff_df, palette='coolwarm')
plt.title("Linear Regression Feature Coefficients")
plt.tight_layout()
plt.show()
```

/tmp/ipykernel_7636/3484898736.py:9: FutureWarning:


Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `y` variable to `hue` and set `legend=False` for the same effect.

```
sns.barplot(x='Coefficient', y='Feature', data=coeff_df, palette='coolwarm')
```



```
[27]: # Residual plot
residuals = y_test - y_pred
plt.scatter(y_pred, residuals)
plt.axhline(y=0, color='r', linestyle='--')
plt.xlabel("Predicted Values")
```

```
plt.ylabel("Residuals")
plt.title("Residual Plot")
plt.show()
```



```
[28]: from sklearn.linear_model import LassoCV

lasso = LassoCV(cv=5, random_state=42)
lasso.fit(X_train, y_train)

# Get features with non-zero coefficients
lasso_features = X_train.columns[lasso.coef_ != 0]
print(" Features selected by Lasso:\n", lasso_features)
```

```
Features selected by Lasso:
Index(['Gender', 'Age', 'Income Stability', 'Location',
      'Loan Amount Request (USD)', 'Expense Type 2', 'Dependents',
      'Credit Score', 'No. of Defaults', 'Has Active Credit Card',
      'Property ID', 'Property Location', 'Co-Applicant'],
      dtype='object')
```

6 Included Plots:

- **Heatmap:** Displays the correlation between features to identify strong positive or negative linear relationships.
- **Boxplot:** Visualizes the spread, median, and potential outliers in numerical variables, useful for detecting skewness or extreme values.
- **Scatter Plot:** Shows the relationship between two numerical variables, helping to identify trends or clusters.
- **Histogram:** Illustrates the frequency distribution of a variable, showing how data points are spread across intervals.

7 Best Practices Followed

- **Consistent Data Preprocessing:** Ensured all features were cleaned, encoded, and standardized uniformly before feeding into the model. This helps in improving model performance and generalizability.
- **Model Validation with Cross-Validation:** Used 5-fold cross-validation to evaluate the model's robustness across different subsets of the data, reducing the chances of overfitting or bias due to a specific data split.

8 Learning Outcomes

- **End-to-End Workflow Understanding:** Gained hands-on experience with the full machine learning pipeline including EDA, preprocessing, feature selection, training, evaluation, and visualization.
- **Interpretation of Model Performance:** Learned how to use statistical metrics like MAE, MSE, RMSE, and R^2 Score to evaluate regression models and interpret residual plots for diagnosing model fit.

GitHub Link

9 Results Table

Table 1: Model Summary: Loan Amount Prediction

Field	Answer
Description Student's Result	Predicting sanctioned loan amount based on applicant's income, credit, and asset details.
Dataset Size (after preprocessing)	30000 records
Train/Test Split Ratio	80:20 (i.e., (<code>test_size=0.2</code>))
Feature(s) Used for Prediction	Age, Income Stability, Loan Amount Request (USD), Dependents, Credit Score, No. of Defaults, Has Active Credit Card, Property Location, Co-Applicant
Model Used	Linear Regression
Cross-Validation Used? (Yes/No)	Yes
If Yes, Number of Folds (K)	5
Reference to CV Results Table	See Table 2
Mean Absolute Error (MAE)	13803.42
Mean Squared Error (MSE)	527445271.77
Root Mean Squared Error (RMSE)	622966.18
R² Score	0.71
Adjusted R² Score	Not calculated
Most Influential Feature(s)	Loan amount request and Income Stability (features with higher positive coefficient)
Observations from Residual Plot	Residuals are fairly spread with slight underestimation for higher values
Interpretation of Predicted vs Actual Plot	Follows an upward trend, but deviations increase with higher amounts
Any Overfitting or Underfitting Observed?	Slight underfitting
Justification	Training and cross-validation scores are similar. There is no strong pattern, but large errors (residuals) on both ends suggest the model isn't fitting complex relationships well.

Table 2: Cross-Validation Results (K = 5)

Fold	MAE	MSE	RMSE	R ² Score
Fold 1	13803.42	527445271.77	22966.18	0.69
Fold 2	13779.90	493351608.13	22211.52	0.70
Fold 3	14030.71	544801753.47	23340.99	0.67
Fold 4	14044.93	513654615.80	22663.95	0.70
Fold 5	13347.16	440761214.18	20994.31	0.73
Average	13801.23	504002892.67	22435.39	0.70