

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
df=pd.read_csv('train.csv')
df.head()
```



| | Customer ID | Name | Gender | Age | Income (USD) | Income Stability | Profession | Type of Employment | Location | Loan Amount Request (USD) | ... | Credit Score | No. of Defaults | Has Active Credit Card | Product |
|---|-------------|-------------------|--------|-----|--------------|------------------|------------|-----------------------|------------|---------------------------|-----|--------------|-----------------|------------------------|---------|
| 0 | C-36995 | Frederica Shealy | F | 56 | 1933.05 | Low | Working | Sales staff | Semi-Urban | 72809.58 | ... | 809.44 | 0 | NaN | |
| 1 | C-33999 | America Calderone | M | 32 | 4952.91 | Low | Working | NaN | Semi-Urban | 46837.47 | ... | 780.40 | 0 | Unpossessed | |
| 2 | C-3770 | Rosetta Verne | F | 65 | 988.19 | High | Pensioner | NaN | Semi-Urban | 45593.04 | ... | 833.15 | 0 | Unpossessed | |
| 3 | C-26480 | Zoe Chitty | F | 65 | NaN | High | Pensioner | NaN | Rural | 80057.92 | ... | 832.70 | 1 | Unpossessed | |
| 4 | C-23459 | Afton Venema | F | 31 | 2614.77 | Low | Working | High skill tech staff | Semi-Urban | 113858.89 | ... | 745.55 | 1 | Active | |

5 rows × 24 columns

```
df.info()
```

[illegible]

| | | | | |
|----|-----------------------------|-------|----------|---------|
| 0 | Customer ID | 30000 | non-null | object |
| 1 | Name | 30000 | non-null | object |
| 2 | Gender | 29947 | non-null | object |
| 3 | Age | 30000 | non-null | int64 |
| 4 | Income (USD) | 25424 | non-null | float64 |
| 5 | Income Stability | 28317 | non-null | object |
| 6 | Profession | 30000 | non-null | object |
| 7 | Type of Employment | 22730 | non-null | object |
| 8 | Location | 30000 | non-null | object |
| 9 | Loan Amount Request (USD) | 30000 | non-null | float64 |
| 10 | Current Loan Expenses (USD) | 29828 | non-null | float64 |
| 11 | Expense Type 1 | 30000 | non-null | object |
| 12 | Expense Type 2 | 30000 | non-null | object |
| 13 | Dependents | 27507 | non-null | float64 |
| 14 | Credit Score | 28297 | non-null | float64 |
| 15 | No. of Defaults | 30000 | non-null | int64 |
| 16 | Has Active Credit Card | 28434 | non-null | object |
| 17 | Property ID | 30000 | non-null | int64 |
| 18 | Property Age | 25150 | non-null | float64 |
| 19 | Property Type | 30000 | non-null | int64 |
| 20 | Property Location | 29644 | non-null | object |
| 21 | Co-Applicant | 30000 | non-null | int64 |
| 22 | Property Price | 30000 | non-null | float64 |
| 23 | Loan Sanction Amount (USD) | 29660 | non-null | float64 |

dtypes: float64(8), int64(5), object(11)
memory usage: 5.5+ MB

```
df['Co-Applicant'].unique()
```

```
array([ 1, 0, -999])
```

```
df['Has Active Credit Card'].unique()
```

```
array([nan, 'Unpossessed', 'Active', 'Inactive'], dtype=object)
```

```
#Removing unnecessary columns
```

```
df = df.drop(columns=["Customer ID", "Name"])
```

```
# Replace -999 with NaN
df['Co-Applicant'] = df['Co-Applicant'].replace(-999, np.nan)

# Option 1: Impute missing values (e.g., assume no co-applicant)
df['Co-Applicant'] = df['Co-Applicant'].fillna(0)

# Fill NaN with 'Unknown'
df['Has Active Credit Card'] = df['Has Active Credit Card'].fillna('Unknown')

# Optional: Encode as ordinal
credit_card_map = {
    'Unpossessed': 0,
    'Inactive': 1,
    'Active': 2,
    'Unknown': -1
}
df['Has Active Credit Card'] = df['Has Active Credit Card'].map(credit_card_map)
```

```
df.isnull().sum()
```



| | |
|-----------------------------|------|
| Gender | 53 |
| Age | 0 |
| Income (USD) | 4576 |
| Income Stability | 1683 |
| Profession | 0 |
| Type of Employment | 7270 |
| Location | 0 |
| Loan Amount Request (USD) | 0 |
| Current Loan Expenses (USD) | 172 |
| Expense Type 1 | 0 |
| Expense Type 2 | 0 |
| Dependents | 2493 |
| Credit Score | 1703 |
| No. of Defaults | 0 |
| Has Active Credit Card | 0 |
| Property ID | 0 |
| Property Age | 4850 |
| Property Type | 0 |
| Property Location | 356 |

```
Co-Applicant                0
Property Price              0
Loan Sanction Amount (USD)  340
dtype: int64
```

```
#Filling null values
df['Gender']=df['Gender'].fillna(df['Gender'].mode()[0])
df['Income (USD)']=df['Income (USD)'].fillna(df['Income (USD)'].median())
df['Income Stability']=df['Income Stability'].fillna(df['Income Stability'].mode()[0])
```

```
#Dropping this column due to presence of more null values and may categories
df['Type of Employment'].unique()
df=df.drop(columns=['Type of Employment'])
```

```
#Current Loan Expenses (USD) - Numeric → fill with median
df['Current Loan Expenses (USD)'] = df['Current Loan Expenses (USD)'].fillna(df['Current Loan Expenses (USD)'].median())
```

```
#Dependents - Numeric → fill with mode (likely a small integer like 1 or 2)
df['Dependents'] = df['Dependents'].fillna(df['Dependents'].mode()[0])
```

```
#Credit Score - Numeric → fill with median
df['Credit Score'] = df['Credit Score'].fillna(df['Credit Score'].median())
```

```
#Property Age - Numeric → fill with median
df['Property Age'] = df['Property Age'].fillna(df['Property Age'].median())
```

```
#Property Location - Categorical → fill with mode
df['Property Location'] = df['Property Location'].fillna(df['Property Location'].mode()[0])
```

```
# Loan Sanction Amount (USD) - Numeric → fill with median
df['Loan Sanction Amount (USD)'] = df['Loan Sanction Amount (USD)'].fillna(df['Loan Sanction Amount (USD)'].median())
```

```
df['Loan Sanction Amount (USD)'] = df['Loan Sanction Amount (USD)'].replace(0, df['Loan Sanction Amount (USD)'].median())
```

```
df.isnull().sum()
```

| | |
|-----------------------------|---|
| Gender | 0 |
| Age | 0 |
| Income (USD) | 0 |
| Income Stability | 0 |
| Profession | 0 |
| Location | 0 |
| Loan Amount Request (USD) | 0 |
| Current Loan Expenses (USD) | 0 |
| Expense Type 1 | 0 |
| Expense Type 2 | 0 |
| Dependents | 0 |
| Credit Score | 0 |
| No. of Defaults | 0 |
| Has Active Credit Card | 0 |
| Property ID | 0 |
| Property Age | 0 |
| Property Type | 0 |
| Property Location | 0 |
| Co-Applicant | 0 |
| Property Price | 0 |
| Loan Sanction Amount (USD) | 0 |
| dtype: int64 | |

Encoding of variables with values

```
from sklearn.preprocessing import LabelEncoder

# List of categorical columns
cat_cols = [
    'Gender', 'Income Stability', 'Profession',
    'Expense Type 1', 'Expense Type 2',
    'Has Active Credit Card', 'Property Type', 'Property Location', 'Location'
]

# Create a label encoder instance
le = LabelEncoder()

# Apply label encoding to each column
for col in cat_cols:
```

```
df[col] = le.fit_transform(df[col])
```

Standardization of Features

```
from sklearn.preprocessing import StandardScaler

# Identify numeric columns (excluding categorical and target)
numeric_cols = df.select_dtypes(include=['float64', 'int64']).columns.tolist()

# Optionally exclude target column (e.g., 'Loan Sanction Amount (USD)')
numeric_cols.remove('Loan Sanction Amount (USD)')

# Initialize scaler
scaler = StandardScaler()

# Fit and transform numeric features
df[numeric_cols] = scaler.fit_transform(df[numeric_cols])
```

```
df.head(10)
```



| | Gender | Age | Income (USD) | Income Stability | Profession | Location | Loan Amount Request (USD) | Current Loan Expenses (USD) | Expense Type 1 | Expense Type 2 | ... | Credit Score | No. of Defaults | Action Crec Ca |
|---|-----------|-----------|--------------|------------------|------------|-----------|---------------------------|-----------------------------|----------------|----------------|-----|--------------|-----------------|----------------|
| 0 | -1.007092 | 0.991451 | -0.061266 | 0.305833 | 0.834973 | 0.142149 | -0.269027 | -0.660358 | -0.749241 | -1.433524 | ... | 0.992493 | -0.490502 | -2.0969 |
| 1 | 0.992958 | -0.504355 | 0.229972 | 0.305833 | 0.834973 | 0.142149 | -0.705269 | 0.392886 | -0.749241 | 0.697582 | ... | 0.578136 | -0.490502 | -1.0017 |
| 2 | -1.007092 | 1.552379 | -0.152389 | -3.269763 | -0.686548 | 0.142149 | -0.726171 | -0.946193 | -0.749241 | 0.697582 | ... | 1.330799 | -0.490502 | -1.0017 |
| 3 | -1.007092 | 1.552379 | -0.033357 | -3.269763 | -0.686548 | -1.762481 | -0.147279 | -0.422775 | -0.749241 | 0.697582 | ... | 1.324379 | 2.038728 | -1.0017 |
| 4 | -1.007092 | -0.566680 | 0.004480 | 0.305833 | 0.834973 | 0.142149 | 0.420461 | 0.374693 | -0.749241 | 0.697582 | ... | 0.080879 | 2.038728 | 1.1885 |
| 5 | -1.007092 | 1.240752 | -0.128594 | 0.305833 | -0.306168 | -1.762481 | -0.913593 | -0.906788 | -0.749241 | -1.433524 | ... | -0.795636 | 2.038728 | 0.0933 |
| 6 | 0.992958 | 0.181223 | -0.019940 | 0.305833 | 0.834973 | 0.142149 | 1.070530 | 1.227526 | 1.334685 | 0.697582 | ... | -1.463830 | -0.490502 | -1.0017 |
| 7 | -1.007092 | 0.305874 | -0.033357 | 0.305833 | -0.306168 | 0.142149 | 2.544436 | 1.682224 | -0.749241 | -1.433524 | ... | 1.032730 | -0.490502 | 1.1885 |
| 8 | -1.007092 | -0.130403 | -0.122697 | 0.305833 | 0.834973 | -1.762481 | -0.901713 | -1.012348 | -0.749241 | 0.697582 | ... | -0.493571 | 2.038728 | 1.1885 |
| 9 | 0.992958 | -1.376908 | -0.098577 | 0.305833 | 0.834973 | -1.762481 | -0.784989 | 0.411038 | -0.749241 | -1.433524 | ... | -1.806987 | -0.490502 | -1.0017 |

10 rows × 21 columns

EDA

Histogram

```
import seaborn as sns
import matplotlib.pyplot as plt


# Plot distributions for selected numeric columns
cols_to_plot = ['Income (USD)', 'Credit Score', 'Loan Amount Request (USD)', 'Loan Sanction Amount (USD)']

for col in cols_to_plot:
    plt.figure(figsize=(6, 4))
```


```
sns.histplot(df[col], kde=True, bins=30)
plt.title(f'Distribution of {col}')
plt.tight_layout()
plt.show()
```



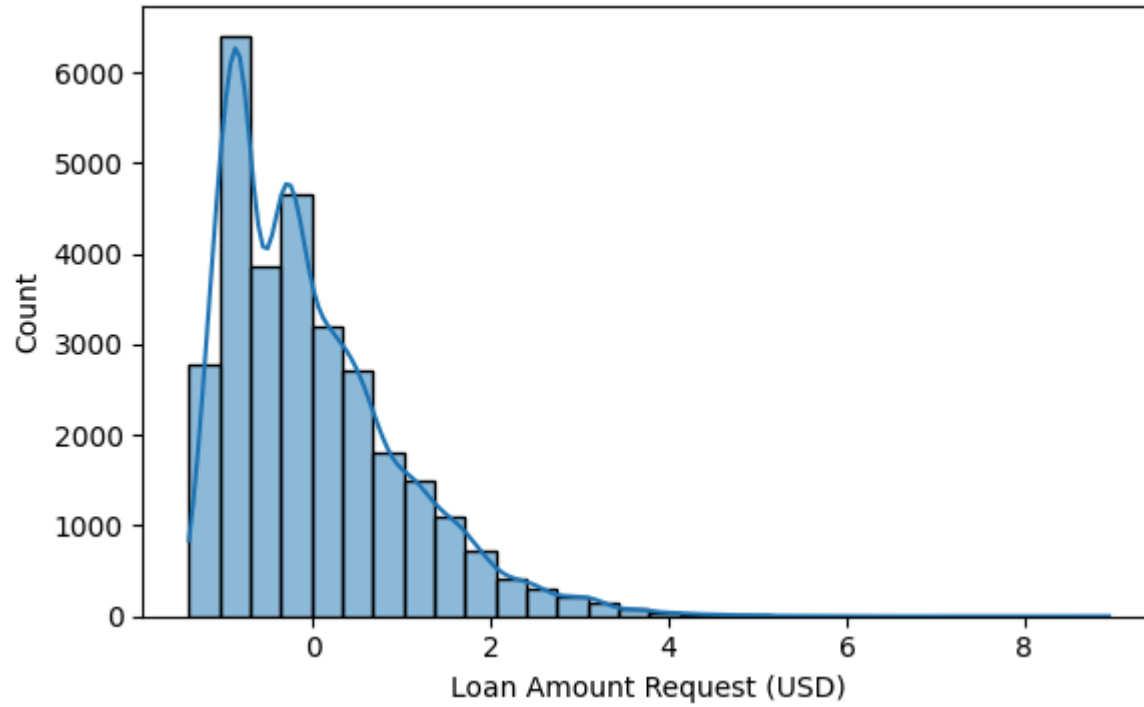

Distribution of Income (USD)




Distribution of Credit Score




Distribution of Loan Amount Request (USD)



Distribution of Loan Sanction Amount (USD)






```
# Income vs Loan Sanction Amount
sns.scatterplot(data=df, x='Income (USD)', y='Loan Sanction Amount (USD)')
plt.title('Income vs Loan Sanction Amount')
plt.show()


# Credit Score vs Loan Sanction Amount
sns.scatterplot(data=df, x='Credit Score', y='Loan Sanction Amount (USD)')
plt.title('Credit Score vs Loan Sanction Amount')
plt.show()
```




Income vs Loan Sanction Amount



Credit Score vs Loan Sanction Amount






Correlation heatmap

```
plt.figure(figsize=(12, 8))
sns.heatmap(df.corr(), annot=True, cmap='coolwarm', fmt=".2f")
plt.title('Correlation Heatmap')
plt.show()
```



Correlation Heatmap



Double-click (or enter) to edit


BoxPlot

```
# Boxplot for Income
sns.boxplot(x=df['Income (USD)'])
plt.title('Boxplot of Income')
plt.show()


# Boxplot for Loan Amount Request
sns.boxplot(x=df['Loan Amount Request (USD)'])
plt.title('Boxplot of Loan Amount Request')
plt.show()
```




Boxplot of Income



Boxplot of Loan Amount Request





Train Test Split

```
from sklearn.model_selection import train_test_split

# Define target variable
target = 'Loan Sanction Amount (USD)'

# Define feature columns
X = df.drop(columns=[target])
y = df[target]

# Split into train and test sets (80% train, 20% test)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

Model Training

```
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score, mean_absolute_error

# Initialize and train the model
model = LinearRegression()
model.fit(X_train, y_train)

# Predict on test set
y_pred = model.predict(X_test)
```

```
# Evaluation Metrics
mse = mean_squared_error(y_test, y_pred)
rmse = np.sqrt(mse)
mae = mean_absolute_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

# Print results
print(f"Mean Squared Error (MSE): {mse:.2f}")
print(f"Root Mean Squared Error (RMSE): {rmse:.2f}")
print(f"Mean Absolute Error (MAE): {mae:.2f}")
print(f"R² Score: {r2:.2f}")
```



```
Mean Squared Error (MSE): 527445271.77
Root Mean Squared Error (RMSE): 22966.18
Mean Absolute Error (MAE): 13803.42
R² Score: 0.69
```

```
from sklearn.model_selection import KFold, cross_val_score
import numpy as np

# Define K-Fold with 5 splits
kf = KFold(n_splits=5, shuffle=True, random_state=42)

# Custom scoring functions
mse_scores = cross_val_score(model, X, y, scoring='neg_mean_squared_error', cv=kf)
mae_scores = cross_val_score(model, X, y, scoring='neg_mean_absolute_error', cv=kf)
r2_scores = cross_val_score(model, X, y, scoring='r2', cv=kf)

# Convert negative MSE/MAE to positive
mse_scores = -mse_scores
mae_scores = -mae_scores
rmse_scores = np.sqrt(mse_scores)

# Print metrics per fold
print("Fold-wise Metrics:")
for i in range(len(mse_scores)):
    print(f"Fold {i+1}:")
```

```
print(f" MSE : {mse_scores[i]:.2f}")
print(f" RMSE: {rmse_scores[i]:.2f}")
print(f" MAE : {mae_scores[i]:.2f}")
print(f" R2 : {r2_scores[i]:.2f}")
print()

# Print average performance
print("Average Metrics Across Folds:")
print(f"Average MSE : {mse_scores.mean():.2f}")
print(f"Average RMSE: {rmse_scores.mean():.2f}")
print(f"Average MAE : {mae_scores.mean():.2f}")
print(f"Average R2 : {r2_scores.mean():.2f}")
```



Fold-wise Metrics:

Fold 1:

```
MSE : 527445271.77
RMSE: 22966.18
MAE : 13803.42
R2 : 0.69
```

Fold 2:

```
MSE : 493351608.13
RMSE: 22211.52
MAE : 13779.90
R2 : 0.70
```

Fold 3:

```
MSE : 544801753.47
RMSE: 23340.99
MAE : 14030.71
R2 : 0.67
```

Fold 4:

```
MSE : 513654615.80
RMSE: 22663.95
MAE : 14044.93
R2 : 0.70
```

Fold 5:

```
MSE : 440761214.18
RMSE: 20994.31
```

MAE : 13347.16

R² : 0.73

Average Metrics Across Folds:

Average MSE : 504002892.67

Average RMSE: 22435.39

Average MAE : 13801.23

Average R² : 0.70

Actual vs Predicted values

```
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.metrics import r2_score

# Predict values
y_pred = model.predict(X_test)

# Plot
plt.figure(figsize=(8, 6))
sns.scatterplot(x=y_test, y=y_pred, alpha=0.7)
plt.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()], color='red', linestyle='--')
plt.xlabel("Actual Loan Sanction Amount")
plt.ylabel("Predicted Loan Sanction Amount")
plt.title(f"Actual vs Predicted (R² = {r2_score(y_test, y_pred):.2f})")
plt.grid(True)
plt.show()
```

