

loan_amount_prediction_svm

September 3, 2025

```
[1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.svm import SVR
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
```

```
[2]: df=pd.read_csv('train.csv')
df.head()
```

```
[2]: Customer ID      Name Gender  Age  Income (USD)  Income Stability \
0      C-36995  Frederica Shealy    F   56      1933.05              Low
1      C-33999  America Calderone    M   32      4952.91              Low
2      C-3770   Rosetta Verne      F   65       988.19              High
3      C-26480    Zoe Chitty      F   65         NaN              High
4      C-23459   Afton Venema      F   31      2614.77              Low
```

```
      Profession      Type of Employment      Location  Loan Amount Request (USD) \
0      Working      Sales staff  Semi-Urban              72809.58
1      Working      NaN  Semi-Urban              46837.47
2  Pensioner      NaN  Semi-Urban              45593.04
3  Pensioner      NaN      Rural              80057.92
4      Working  High skill tech staff  Semi-Urban              113858.89
```

```
      ...  Credit Score  No. of Defaults  Has Active Credit Card  Property ID \
0      ...      809.44              0              NaN              746
1      ...      780.40              0      Unpossessed              608
2      ...      833.15              0      Unpossessed              546
3      ...      832.70              1      Unpossessed              890
4      ...      745.55              1      Active              715
```

```
      Property Age  Property Type  Property Location  Co-Applicant \
0      1933.05      4      Rural      1
1      4952.91      2      Rural      1
2      988.19      2      Urban      0
3      NaN      2      Semi-Urban      1
4      2614.77      4      Semi-Urban      1
```

	Property Price	Loan Sanction Amount (USD)
0	119933.46	54607.18
1	54791.00	37469.98
2	72440.58	36474.43
3	121441.51	56040.54
4	208567.91	74008.28

[5 rows x 24 columns]

```
[3]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 30000 entries, 0 to 29999
Data columns (total 24 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Customer ID                          30000 non-null  object
1   Name                                30000 non-null  object
2   Gender                              29947 non-null  object
3   Age                                 30000 non-null  int64
4   Income (USD)                        25424 non-null  float64
5   Income Stability                     28317 non-null  object
6   Profession                           30000 non-null  object
7   Type of Employment                  22730 non-null  object
8   Location                            30000 non-null  object
9   Loan Amount Request (USD)           30000 non-null  float64
10  Current Loan Expenses (USD)         29828 non-null  float64
11  Expense Type 1                      30000 non-null  object
12  Expense Type 2                      30000 non-null  object
13  Dependents                          27507 non-null  float64
14  Credit Score                        28297 non-null  float64
15  No. of Defaults                     30000 non-null  int64
16  Has Active Credit Card              28434 non-null  object
17  Property ID                         30000 non-null  int64
18  Property Age                        25150 non-null  float64
19  Property Type                       30000 non-null  int64
20  Property Location                   29644 non-null  object
21  Co-Applicant                       30000 non-null  int64
22  Property Price                      30000 non-null  float64
23  Loan Sanction Amount (USD)          29660 non-null  float64
dtypes: float64(8), int64(5), object(11)
memory usage: 5.5+ MB
```

```
[4]: df['Co-Applicant'].unique()
```

```
[4]: array([ 1, 0, -999])
```

```
[5]: df['Has Active Credit Card'].unique()
```

```
[5]: array([nan, 'Unpossessed', 'Active', 'Inactive'], dtype=object)
```

```
[6]: #Removing unnecessary columns  
df = df.drop(columns=["Customer ID", "Name"])
```

```
[7]: # Replace -999 with NaN  
df['Co-Applicant'] = df['Co-Applicant'].replace(-999, np.nan)  
  
# Option 1: Impute missing values (e.g., assume no co-applicant)  
df['Co-Applicant'] = df['Co-Applicant'].fillna(0)
```

```
[8]: # Fill NaN with 'Unknown'  
df['Has Active Credit Card'] = df['Has Active Credit Card'].fillna('Unknown')  
  
# Optional: Encode as ordinal  
credit_card_map = {  
    'Unpossessed': 0,  
    'Inactive': 1,  
    'Active': 2,  
    'Unknown': -1  
}  
df['Has Active Credit Card'] = df['Has Active Credit Card'].map(credit_card_map)
```

```
[9]: df.isnull().sum()
```

```
[9]: Gender                53  
Age                      0  
Income (USD)            4576  
Income Stability        1683  
Profession              0  
Type of Employment     7270  
Location                0  
Loan Amount Request (USD) 0  
Current Loan Expenses (USD) 172  
Expense Type 1          0  
Expense Type 2          0  
Dependents              2493  
Credit Score           1703  
No. of Defaults         0  
Has Active Credit Card  0  
Property ID            0  
Property Age           4850  
Property Type          0  
Property Location      356  
Co-Applicant            0
```

```
Property Price          0
Loan Sanction Amount (USD) 340
dtype: int64
```

```
[10]: #Filling null values
df['Gender']=df['Gender'].fillna(df['Gender'].mode()[0])
df['Income (USD)']=df['Income (USD)'].fillna(df['Income (USD)'].median())
df['Income Stability']=df['Income Stability'].fillna(df['Income Stability'].
↳mode()[0])
```

```
[11]: #Dropping this column due to presence of more null values and may categories
df['Type of Employment'].unique()
df=df.drop(columns=['Type of Employment'])
```

```
[12]: #Current Loan Expenses (USD) - Numeric → fill with median
df['Current Loan Expenses (USD)'] = df['Current Loan Expenses (USD)'].
↳fillna(df['Current Loan Expenses (USD)'].median())

#Dependents - Numeric → fill with mode (likely a small integer like 1 or 2)
df['Dependents'] = df['Dependents'].fillna(df['Dependents'].mode()[0])

#Credit Score - Numeric → fill with median
df['Credit Score'] = df['Credit Score'].fillna(df['Credit Score'].median())

#Property Age - Numeric → fill with median
df['Property Age'] = df['Property Age'].fillna(df['Property Age'].median())

#Property Location - Categorical → fill with mode
df['Property Location'] = df['Property Location'].fillna(df['Property_
↳Location'].mode()[0])

# Loan Sanction Amount (USD) - Numeric → fill with median
df['Loan Sanction Amount (USD)'] = df['Loan Sanction Amount (USD)'].
↳fillna(df['Loan Sanction Amount (USD)'].median())

df['Loan Sanction Amount (USD)'] = df['Loan Sanction Amount (USD)'].replace(0,↳
↳df['Loan Sanction Amount (USD)'].median())
```

```
[13]: df.isnull().sum()
```

```
[13]: Gender          0
Age                  0
Income (USD)         0
Income Stability     0
Profession           0
Location             0
Loan Amount Request (USD) 0
```

Current Loan Expenses (USD)	0
Expense Type 1	0
Expense Type 2	0
Dependents	0
Credit Score	0
No. of Defaults	0
Has Active Credit Card	0
Property ID	0
Property Age	0
Property Type	0
Property Location	0
Co-Applicant	0
Property Price	0
Loan Sanction Amount (USD)	0

dtype: int64

Encoding of variables with values

```
[14]: from sklearn.preprocessing import LabelEncoder

# List of categorical columns
cat_cols = [
    'Gender', 'Income Stability', 'Profession',
    'Expense Type 1', 'Expense Type 2',
    'Has Active Credit Card', 'Property Type', 'Property Location', 'Location'
]

# Create a label encoder instance
le = LabelEncoder()

# Apply label encoding to each column
for col in cat_cols:
    df[col] = le.fit_transform(df[col])
```

Standardization of Features

```
[15]: from sklearn.preprocessing import StandardScaler

# Identify numeric columns (excluding categorical and target)
numeric_cols = df.select_dtypes(include=['float64', 'int64']).columns.tolist()

# Optionally exclude target column (e.g., 'Loan Sanction Amount (USD)')
numeric_cols.remove('Loan Sanction Amount (USD)')

# Initialize scaler
scaler = StandardScaler()

# Fit and transform numeric features
```

```
df[numeric_cols] = scaler.fit_transform(df[numeric_cols])
```

```
[16]: df.head(10)
```

```
[16]:      Gender      Age  Income (USD)  Income Stability  Profession  Location \
0 -1.007092  0.991451   -0.061266      0.305833    0.834973  0.142149
1  0.992958 -0.504355    0.229972      0.305833    0.834973  0.142149
2 -1.007092  1.552379   -0.152389     -3.269763   -0.686548  0.142149
3 -1.007092  1.552379   -0.033357     -3.269763   -0.686548 -1.762481
4 -1.007092 -0.566680    0.004480      0.305833    0.834973  0.142149
5 -1.007092  1.240752   -0.128594      0.305833   -0.306168 -1.762481
6  0.992958  0.181223   -0.019940      0.305833    0.834973  0.142149
7 -1.007092  0.305874   -0.033357      0.305833   -0.306168  0.142149
8 -1.007092 -0.130403   -0.122697      0.305833    0.834973 -1.762481
9  0.992958 -1.376908   -0.098577      0.305833    0.834973 -1.762481
```

```
      Loan Amount Request (USD)  Current Loan Expenses (USD)  Expense Type 1 \
0                -0.269027                -0.660358                -0.749241
1                -0.705269                 0.392886                -0.749241
2                -0.726171                -0.946193                -0.749241
3                -0.147279                -0.422775                -0.749241
4                 0.420461                 0.374693                -0.749241
5                -0.913593                -0.906788                -0.749241
6                 1.070530                 1.227526                 1.334685
7                 2.544436                 1.682224                -0.749241
8                -0.901713                -1.012348                -0.749241
9                -0.784989                 0.411038                -0.749241
```

```
      Expense Type 2 ... Credit Score  No. of Defaults  Has Active Credit Card \
0      -1.433524 ...      0.992493      -0.490502                -2.096903
1       0.697582 ...      0.578136      -0.490502                -1.001762
2       0.697582 ...      1.330799      -0.490502                -1.001762
3       0.697582 ...      1.324379       2.038728                -1.001762
4       0.697582 ...      0.080879       2.038728                 1.188520
5      -1.433524 ...     -0.795636       2.038728                 0.093379
6       0.697582 ...     -1.463830      -0.490502                -1.001762
7      -1.433524 ...      1.032730      -0.490502                 1.188520
8       0.697582 ...     -0.493571       2.038728                 1.188520
9      -1.433524 ...     -1.806987      -0.490502                -1.001762
```

```
      Property ID  Property Age  Property Type  Property Location  Co-Applicant \
0      0.846998   -0.060969      1.376731      -1.214540      0.419205
1      0.368086    0.230298     -0.411309     -1.214540      0.419205
2      0.152923   -0.152102     -0.411309      1.283229     -2.385467
3      1.346732   -0.032979     -0.411309      0.034344      0.419205
4      0.739417    0.004783      1.376731      0.034344      0.419205
5     -0.037948   -0.128305     -0.411309     -1.214540      0.419205
```

6	-0.954127	-0.019639	-1.305329	0.034344	0.419205
7	-0.652204	-0.032979	-0.411309	1.283229	0.419205
8	-0.905541	-0.122407	1.376731	-1.214540	0.419205
9	1.322440	-0.098284	-0.411309	1.283229	0.419205

	Property Price	Loan Sanction Amount (USD)
0	-0.126419	54607.180
1	-0.822772	37469.980
2	-0.634103	36474.430
3	-0.110298	56040.540
4	0.821057	74008.280
5	-0.947245	22382.570
6	0.954495	35209.395
7	2.878533	168218.240
8	-0.821570	22842.290
9	-0.681642	35209.395

[10 rows x 21 columns]



EDA



Histogram

```
[17]: import seaborn as sns
import matplotlib.pyplot as plt

# Plot distributions for selected numeric columns
cols_to_plot = ['Income (USD)', 'Credit Score', 'Loan Amount Request (USD)', 'Loan Sanction Amount (USD)']

for col in cols_to_plot:
    plt.figure(figsize=(6, 4))
    sns.histplot(df[col], kde=True, bins=30)
    plt.title(f'Distribution of {col}')
    plt.tight_layout()
    plt.show()
```







Scatter Plot

```
[18]: # Income vs Loan Sanction Amount
sns.scatterplot(data=df, x='Income (USD)', y='Loan Sanction Amount (USD)')
plt.title('Income vs Loan Sanction Amount')
plt.show()


# Credit Score vs Loan Sanction Amount
sns.scatterplot(data=df, x='Credit Score', y='Loan Sanction Amount (USD)')
plt.title('Credit Score vs Loan Sanction Amount')
plt.show()
```





Correlation heatmap

```
[19]: plt.figure(figsize=(12, 8))
sns.heatmap(df.corr(), annot=True, cmap='coolwarm', fmt=".2f")
plt.title('Correlation Heatmap')
plt.show()
```




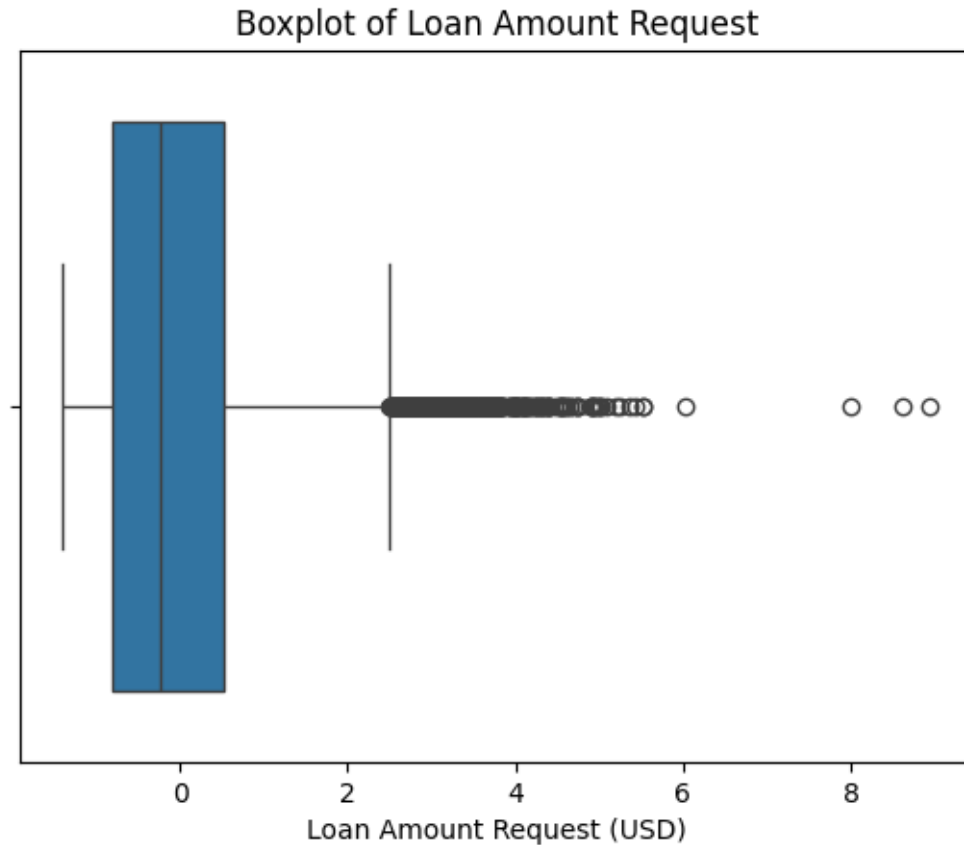
BoxPlot

```
[20]: # Boxplot for Income
sns.boxplot(x=df['Income (USD)'])
plt.title('Boxplot of Income')
plt.show()

# Boxplot for Loan Amount Request
sns.boxplot(x=df['Loan Amount Request (USD)'])
plt.title('Boxplot of Loan Amount Request')
plt.show()
```

Boxplot of Income





```
[21]: # Define target variable
target = 'Loan Sanction Amount (USD)'

# Define feature columns
X = df.drop(columns=[target])
y = df[target]

# Split into train and test sets (80% train, 20% test)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↪random_state=42)
```

0.0.1 Grid Search

```
[23]: # Define the parameter grid
param_grid = {
    'kernel': ['linear', 'rbf'],
    'C': [0.1, 1, 10],
    'gamma': ['scale', 0.1]
}
```

```

# Split into train (80%) and test (20%)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳random_state=42)

# Initialize SVR model
svr = SVR()

# Grid Search with 5-fold cross-validation
grid_search = GridSearchCV(
    estimator=svr,
    param_grid=param_grid,
    scoring='r2',
    cv=5,
    n_jobs=-1,
    verbose=0
)

# Fit on training set
grid_search.fit(X_train, y_train)

# Best parameters and score
print("Best Parameters:", grid_search.best_params_)
print("Best CV Score (R²):", grid_search.best_score_)

```

Best Parameters: {'C': 10, 'gamma': 'scale', 'kernel': 'linear'}

Best CV Score (R²): 0.6543903124032083

```

[25]: # Define target variable
target = 'Loan Sanction Amount (USD)'

# Define feature columns
X = df.drop(columns=[target])
y = df[target]

# Split into train and test sets (80% train, 20% test)
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42
)

# Train SVR with best parameters on the training set
best_svr = grid_search.best_estimator_
best_svr.fit(X_train, y_train)

# Predictions on test set
y_pred = best_svr.predict(X_test)

# Evaluation metrics

```

```
mse = mean_squared_error(y_test, y_pred)
mae = mean_absolute_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

print("\nModel Evaluation on Test Set:")
print(f"Mean Squared Error (MSE): {mse:.4f}")
print(f"Mean Absolute Error (MAE): {mae:.4f}")
print(f"R2 Score: {r2:.4f}")
```

```
Model Evaluation on Test Set:
Mean Squared Error (MSE): 552578763.6120
Mean Absolute Error (MAE): 11951.5667
R2 Score: 0.6708
```