

mlp

September 9, 2025

```
[12]: import os
import numpy as np
import pandas as pd
from PIL import Image
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder, label_binarize
from sklearn.metrics import (
    classification_report, confusion_matrix,
    accuracy_score, roc_curve, auc
)

import torch
import torch.nn as nn
import torch.optim as optim
from torch.utils.data import DataLoader, TensorDataset

import matplotlib.pyplot as plt
import seaborn as sns
```

```
[13]: csv_path = "archive/english.csv"
df = pd.read_csv(csv_path)

ROOT = "archive"
IMG_SIZE = (28, 28) # width, height

images, labels, missing = [], [], []
for idx, row in df.iterrows():
    img_rel = row["image"]
    label_raw = row["label"]
    img_path = os.path.join(ROOT, img_rel)

    if not os.path.exists(img_path):
        missing.append(img_path)
        continue

    im = Image.open(img_path).convert("L")
    im = im.resize(IMG_SIZE)
    arr = np.asarray(im, dtype=np.float32).flatten() / 255.0 # normalize [0,1]
```

```

        images.append(arr)
        labels.append(label_raw)

print("Missing files:", len(missing))

X = np.vstack(images) # (N, D)
y_raw = np.array(labels)
print("Loaded X:", X.shape, " y:", y_raw.shape)

```

Missing files: 0
Loaded X: (3410, 784) y: (3410,)

```

[14]: # Encode to integers 0..C-1
le = LabelEncoder()
y = le.fit_transform(y_raw)

# Stratified split -> train, validation, test
RNG_SEED = 42
X_train_full, X_test, y_train_full, y_test = train_test_split(
    X, y, test_size=0.2, stratify=y, random_state=RNG_SEED
)

X_train, X_val, y_train, y_val = train_test_split(
    X_train_full, y_train_full, test_size=0.2,
    stratify=y_train_full, random_state=RNG_SEED
)

print("Shapes -> train:", X_train.shape, "val:", X_val.shape, "test:", X_test.
      ↪shape)
print("Num classes:", len(le.classes_))

```

Shapes -> train: (2182, 784) val: (546, 784) test: (682, 784)
Num classes: 62

```

[15]: X_train_t = torch.tensor(X_train, dtype=torch.float32)
y_train_t = torch.tensor(y_train, dtype=torch.long)
X_val_t = torch.tensor(X_val, dtype=torch.float32)
y_val_t = torch.tensor(y_val, dtype=torch.long)
X_test_t = torch.tensor(X_test, dtype=torch.float32)
y_test_t = torch.tensor(y_test, dtype=torch.long)

def get_loader(X, y, batch_size):
    ds = TensorDataset(X, y)
    return DataLoader(ds, batch_size=batch_size, shuffle=True)

```

```

[16]: class MLP(nn.Module):

```

```

def __init__(self, input_dim, hidden_dim, output_dim, activation="relu",
↳num_hidden=1):
    super().__init__()
    act_fn = {"relu": nn.ReLU(), "tanh": nn.Tanh(), "sigmoid": nn.
↳Sigmoid()}[activation]

    layers = []
    layers.append(nn.Linear(input_dim, hidden_dim))
    layers.append(act_fn)

    if num_hidden == 2:
        layers.append(nn.Linear(hidden_dim, hidden_dim))
        layers.append(act_fn)

    layers.append(nn.Linear(hidden_dim, output_dim))
    self.net = nn.Sequential(*layers)

def forward(self, x):
    return self.net(x)

```

```

[17]: def train_model(params):
    batch_size, lr, hidden_dim, activation, optimizer_name, num_hidden = params

    train_loader = get_loader(X_train_t, y_train_t, batch_size)
    val_loader = get_loader(X_val_t, y_val_t, batch_size)

    model = MLP(X_train.shape[1], hidden_dim, len(le.classes_), activation,
↳num_hidden)
    criterion = nn.CrossEntropyLoss()

    if optimizer_name == "sgd":
        optimizer = optim.SGD(model.parameters(), lr=lr)
    else:
        optimizer = optim.Adam(model.parameters(), lr=lr)

    history = {"train_loss": [], "val_loss": [], "val_acc": []}
    EPOCHS = 20

    for epoch in range(EPOCHS):
        model.train()
        train_loss = 0
        for xb, yb in train_loader:
            optimizer.zero_grad()
            out = model(xb)
            loss = criterion(out, yb)
            loss.backward()
            optimizer.step()

```

```

        train_loss += loss.item()

    # validation
    model.eval()
    val_loss, correct = 0, 0
    with torch.no_grad():
        for xb, yb in val_loader:
            out = model(xb)
            loss = criterion(out, yb)
            val_loss += loss.item()
            preds = out.argmax(dim=1)
            correct += (preds == yb).sum().item()
    acc = correct / len(val_loader.dataset)

    history["train_loss"].append(train_loss/len(train_loader))
    history["val_loss"].append(val_loss/len(val_loader))
    history["val_acc"].append(acc)

    return model, history

```

```

[11]: search_space = [
    (bs, lr, hd, act, opt, nh)
    for bs in [32, 64, 128]           # batch sizes
    for lr in [0.1, 0.01, 0.001]      # learning rates
    for hd in [128, 256]              # hidden layer size
    for act in ["relu", "tanh", "sigmoid"] # activations
    for opt in ["sgd", "adam"]         # optimizers
    for nh in [1, 2]                  # number of hidden layers
]

best_acc, best_params, best_model, best_history = 0, None, None, None

for params in search_space:
    model, hist = train_model(params)
    final_acc = hist["val_acc"][-1]

    if final_acc > best_acc:
        best_acc = final_acc
        best_params = params
        best_model = model
        best_history = hist

print("\nBest Params:", best_params)
print("Best Val Accuracy:", best_acc)

```

```

Best Params: (32, 0.001, 256, 'sigmoid', 'adam', 1)
Best Val Accuracy: 0.30036630036630035

```

```
[18]: best_model.eval()
with torch.no_grad():
    preds = best_model(X_test_t).argmax(dim=1).numpy()

print("\nMLP Test Accuracy:", accuracy_score(y_test, preds))
print("\nClassification Report:\n", classification_report(y_test, preds,
    ↪target_names=le.classes_))

# Confusion Matrix
cm = confusion_matrix(y_test, preds)
plt.figure(figsize=(10,8))
sns.heatmap(cm, cmap="Blues", xticklabels=le.classes_, yticklabels=le.classes_,
    ↪annot=False)
plt.xlabel("Predicted")
plt.ylabel("True")
plt.title("MLP Confusion Matrix")
plt.show()
```

MLP Test Accuracy: 0.2668621700879765

Classification Report:

	precision	recall	f1-score	support
0	0.00	0.00	0.00	11
1	0.15	0.18	0.17	11
2	0.20	0.09	0.12	11
3	0.20	0.09	0.12	11
4	0.00	0.00	0.00	11
5	0.00	0.00	0.00	11
6	0.40	0.18	0.25	11
7	0.19	0.27	0.22	11
8	0.44	0.73	0.55	11
9	0.23	0.45	0.30	11
A	0.32	0.55	0.40	11
B	1.00	0.09	0.17	11
C	0.54	0.64	0.58	11
D	1.00	0.09	0.17	11
E	1.00	0.09	0.17	11
F	0.20	0.18	0.19	11
G	0.31	0.36	0.33	11
H	0.67	0.18	0.29	11
I	0.30	0.55	0.39	11
J	0.33	0.09	0.14	11
K	0.22	0.18	0.20	11
L	0.35	0.82	0.49	11
M	0.37	0.64	0.47	11
N	0.33	0.18	0.24	11

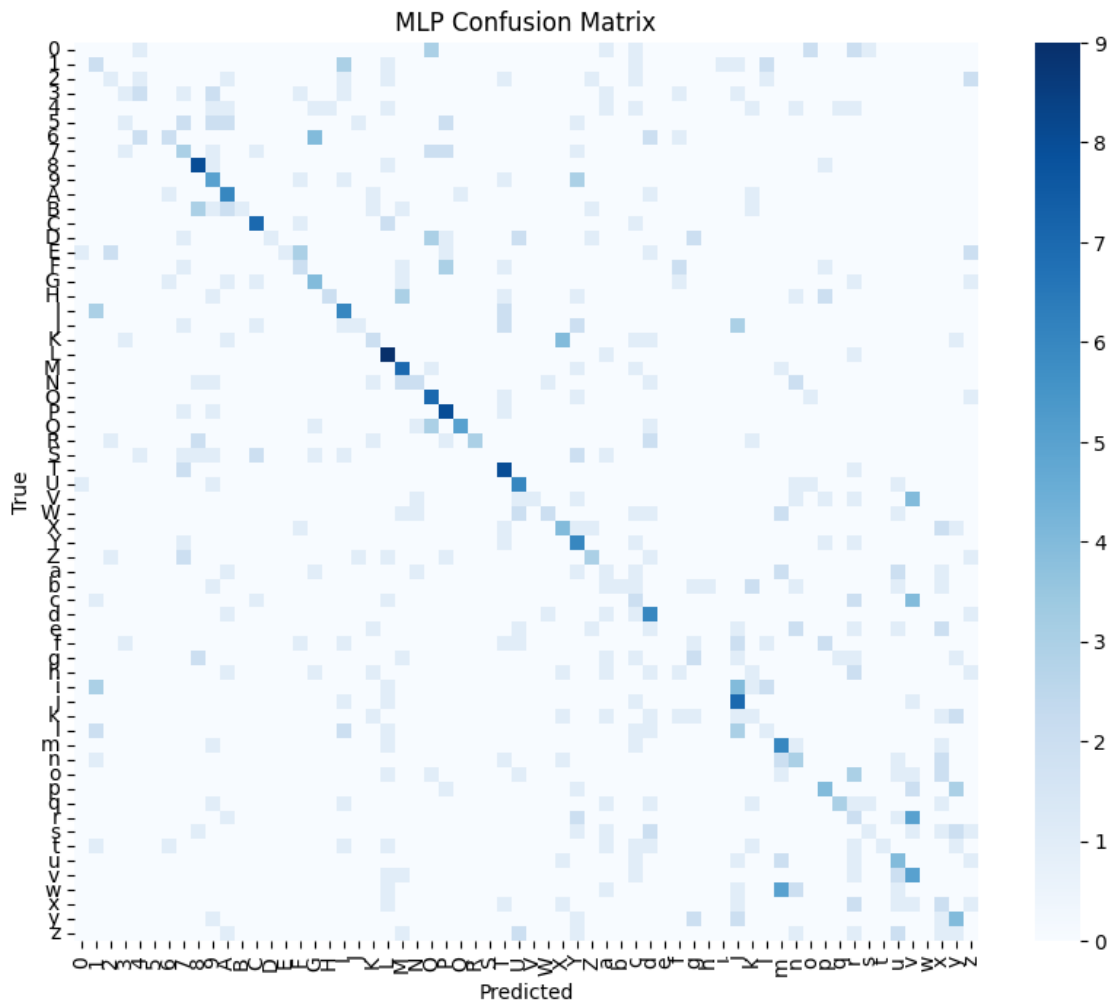
O	0.35	0.64	0.45	11
P	0.40	0.73	0.52	11
Q	0.83	0.45	0.59	11
R	1.00	0.27	0.43	11
S	0.00	0.00	0.00	11
T	0.35	0.73	0.47	11
U	0.38	0.55	0.44	11
V	1.00	0.09	0.17	11
W	0.50	0.18	0.27	11
X	0.31	0.36	0.33	11
Y	0.22	0.55	0.32	11
Z	0.38	0.27	0.32	11
a	0.07	0.09	0.08	11
b	1.00	0.09	0.17	11
c	0.10	0.18	0.12	11
d	0.24	0.55	0.33	11
e	0.00	0.00	0.00	11
f	0.00	0.00	0.00	11
g	0.22	0.18	0.20	11
h	0.00	0.00	0.00	11
i	0.00	0.00	0.00	11
j	0.24	0.64	0.35	11
k	0.09	0.09	0.09	11
l	0.14	0.09	0.11	11
m	0.29	0.55	0.38	11
n	0.19	0.27	0.22	11
o	0.00	0.00	0.00	11
p	0.33	0.36	0.35	11
q	0.60	0.27	0.38	11
r	0.08	0.18	0.11	11
s	0.33	0.09	0.14	11
t	1.00	0.09	0.17	11
u	0.21	0.36	0.27	11
v	0.21	0.45	0.29	11
w	0.00	0.00	0.00	11
x	0.11	0.18	0.14	11
y	0.25	0.36	0.30	11
z	0.00	0.00	0.00	11
accuracy			0.27	682
macro avg	0.33	0.27	0.23	682
weighted avg	0.33	0.27	0.23	682

/home/pranesh/Downloads/ML Lab/.venv/lib/python3.12/site-packages/sklearn/metrics/_classification.py:1731: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.

```

_warn_prf(average, modifier, f"{metric.capitalize()} is", result.shape[0])
/home/pranesh/Downloads/ML Lab/.venv/lib/python3.12/site-
packages/sklearn/metrics/_classification.py:1731: UndefinedMetricWarning:
Precision is ill-defined and being set to 0.0 in labels with no predicted
samples. Use `zero_division` parameter to control this behavior.
_warn_prf(average, modifier, f"{metric.capitalize()} is", result.shape[0])
/home/pranesh/Downloads/ML Lab/.venv/lib/python3.12/site-
packages/sklearn/metrics/_classification.py:1731: UndefinedMetricWarning:
Precision is ill-defined and being set to 0.0 in labels with no predicted
samples. Use `zero_division` parameter to control this behavior.
_warn_prf(average, modifier, f"{metric.capitalize()} is", result.shape[0])

```



```

[19]: y_test_bin = label_binarize(y_test, classes=np.arange(len(le.classes_)))
      preds_bin = label_binarize(preds, classes=np.arange(len(le.classes_)))

      fpr_micro, tpr_micro, _ = roc_curve(y_test_bin.ravel(), preds_bin.ravel())

```

```


roc_auc_micro = auc(fpr_micro, tpr_micro)

fpr_dict, tpr_dict, roc_auc_dict = {}, {}, {}
for i in range(len(le.classes_)):
    fpr_dict[i], tpr_dict[i], _ = roc_curve(y_test_bin[:, i], preds_bin[:, i])
    roc_auc_dict[i] = auc(fpr_dict[i], tpr_dict[i])

all_fpr = np.unique(np.concatenate([fpr_dict[i] for i in range(len(le.
    ↪classes_))]))
mean_tpr = np.zeros_like(all_fpr)
for i in range(len(le.classes_)):
    mean_tpr += np.interp(all_fpr, fpr_dict[i], tpr_dict[i])
mean_tpr /= len(le.classes_)
roc_auc_macro = auc(all_fpr, mean_tpr)

plt.figure()
plt.plot(fpr_micro, tpr_micro, label=f"Micro-average ROC (AUC = {roc_auc_micro:.
    ↪2f})", color="blue")
plt.plot(all_fpr, mean_tpr, label=f"Macro-average ROC (AUC = {roc_auc_macro:.
    ↪2f})", color="red")
plt.plot([0, 1], [0, 1], "k--")
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("MLP ROC Curves (Micro & Macro)")
plt.legend(loc="lower right")
plt.show()

```

```
[20]: plt.figure()
plt.plot(best_history["train_loss"], label="Train Loss")
plt.plot(best_history["val_loss"], label="Validation Loss")
plt.xlabel("Epochs")
plt.ylabel("Loss")
plt.title("MLP Loss vs Epochs")
plt.legend()
plt.show()

plt.figure()
plt.plot(best_history["val_acc"], label="Validation Accuracy")
plt.xlabel("Epochs")
plt.ylabel("Accuracy")
plt.title("MLP Accuracy vs Epochs")
plt.legend()
plt.show()
```

