

# **Sri Sivasubramaniya Nadar College of Engineering, Chennai**

(An autonomous Institution affiliated to Anna University)

**Degree & Branch:** B.E. Computer Science & Engineering

**Semester:** V

**Subject Code & Name:** ICS1512 - Machine Learning  
Algorithms Laboratory

**Academic Year:** 2025–2026 (Odd)

**Batch:** 2023–2028

**Experiment 6:** Dimensionality Reduction and Model Evaluation (With and Without PCA)

## **Aim**

To study the effect of dimensionality reduction using Principal Component Analysis (PCA) on the performance of various machine learning classifiers. The task requires:

1. Training and validating models without PCA (original feature space).
2. Training and validating models with PCA (reduced feature space).

For both cases, students must perform hyperparameter tuning, apply 5-fold cross-validation, and record performance.

## **Libraries Used**

- Pandas
- NumPy
- Matplotlib
- Scikit-learn
- XGBoost

## **Objective**

To evaluate how dimensionality reduction using Principal Component Analysis (PCA) influences the accuracy and generalization of different machine learning classifiers, by comparing their performance with and without PCA through hyperparameter tuning and 5-fold cross-validation. [a4paper,12pt]article pdfpages

## Including PDF

Here is my PDF included below:

## Experiment 6: Dimensionality Reduction and Model Evaluation (With and Without PCA)

```
# IMPORTS
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import GridSearchCV, RandomizedSearchCV, cross_val_score
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.linear_model import LogisticRegression
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.metrics import f1_score, accuracy_score, make_scorer, roc_curve,
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import AdaBoostClassifier, GradientBoostingClassifier, R
from xgboost import XGBClassifier
from sklearn.svm import SVC
from sklearn.naive_bayes import GaussianNB
from sklearn.neighbors import KNeighborsClassifier
from sklearn.decomposition import PCA
from sklearn.pipeline import Pipeline
```

```
TARGET_COLUMN = 'class'
```

```
df = pd.read_csv('/content/drive/MyDrive/ml-lab/spambase_csv.csv')
print(df.head())
print("Initial shape:", df.shape)
```

	word_freq_make	word_freq_address	word_freq_all	word_freq_3d	\
0	0.00	0.64	0.64	0.0	
1	0.21	0.28	0.50	0.0	
2	0.06	0.00	0.71	0.0	
3	0.00	0.00	0.00	0.0	
4	0.00	0.00	0.00	0.0	

	word_freq_our	word_freq_over	word_freq_remove	word_freq_internet	\
0	0.32	0.00	0.00	0.00	
1	0.14	0.28	0.21	0.07	
2	1.23	0.19	0.19	0.12	
3	0.63	0.00	0.31	0.63	
4	0.63	0.00	0.31	0.63	

	word_freq_order	word_freq_mail	...	char_freq_%3B	char_freq_%28	\
0	0.00	0.00	...	0.00	0.000	
1	0.00	0.94	...	0.00	0.132	
2	0.64	0.25	...	0.01	0.143	
3	0.31	0.63	...	0.00	0.137	
4	0.31	0.63	...	0.00	0.135	

	char_freq_%5B	char_freq_%21	char_freq_%24	char_freq_%23	\
0	0.0	0.778	0.000	0.000	
1	0.0	0.372	0.180	0.048	
2	0.0	0.276	0.184	0.010	
3	0.0	0.137	0.000	0.000	
4	0.0	0.135	0.000	0.000	

	capital_run_length_average	capital_run_length_longest	\
0	3.756	61	
1	5.114	101	
2	9.821	485	
3	3.537	40	
4	3.537	40	

	capital_run_length_total	class
0	278	1
1	1028	1
2	2259	1
3	191	1
4	191	1

[5 rows x 58 columns]  
Initial shape: (4601, 58)

```
# HANDLE MISSING VALUES
df = df.dropna(thresh=df.shape[1]//2) # Drop rows with >50% missing
df.fillna(df.median(numeric_only=True), inplace=True)
```

```
# OUTLIER HANDLING (Z-Score)
def remove_outliers(df, threshold=3):
    numeric_cols = df.select_dtypes(include=[np.number]).columns
    z_scores = np.abs((df[numeric_cols] - df[numeric_cols].mean()) / df[numeric_cols].std())
    return df[(z_scores < threshold).all(axis=1)]

df = remove_outliers(df)
# print("After outlier removal:", df.shape)
```

```
# FEATURE / TARGET SPLIT
X = df.drop(columns=[TARGET_COLUMN])
y = df[TARGET_COLUMN]

# ENCODE + STANDARDIZE
numeric_cols = X.select_dtypes(include=[np.number]).columns
categorical_cols = X.select_dtypes(exclude=[np.number]).columns
X_encoded = pd.get_dummies(X, columns=categorical_cols)
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X_encoded)

# TRAIN / TEST SPLIT
X_train, X_test, y_train, y_test = train_test_split(
    X_scaled,
    y,
    test_size=0.2,          # 20% test hold-out
    stratify=y,             # keep class balance if classification
    random_state=42
)
```

## PCA Variance Explained

```
pca = PCA(n_components=0.95)
pca.fit(X_scaled)

print("Chosen components:", pca.n_components_)
print("Total variance explained (%):", pca.explained_variance_ratio_.sum()*100)
```

Chosen components: 49  
Total variance explained (%): 95.53617010131482

## Support Vector Machine (SVM)

```
# ===== HYPERPARAM GRID =====
param_grid = {
    'kernel': ['linear', 'rbf'],
    'C': [0.1, 10],
    'gamma': ['scale']
}

# ===== HELPER FUNCTION =====
def evaluate_svc(X_train, X_test, y_train, y_test, use_pca=False, pca_variance=0.95):
    if use_pca:
        pca = PCA(n_components=pca_variance)
        X_train_proc = pca.fit_transform(X_train)
        X_test_proc = pca.transform(X_test)
    else:
        X_train_proc = X_train
        X_test_proc = X_test

    svc = SVC(probability=True)
    grid = GridSearchCV(svc, param_grid, cv=5, scoring='accuracy') # you can switch scoring
    grid.fit(X_train_proc, y_train)

    best_model = grid.best_estimator_
    y_pred = best_model.predict(X_test_proc)
    y_proba = best_model.predict_proba(X_test_proc)[:,-1]

    acc = accuracy_score(y_test, y_pred)
    auc = roc_auc_score(y_test, y_proba)

    return grid.best_params_, acc, auc, y_test, y_proba

# ===== EVALUATE ALL COMBOS =====
results = []
for kernel in param_grid['kernel']:
```

```

for C in param_grid['C']:
    for gamma in param_grid['gamma']:
        params = {'kernel': kernel, 'C': C, 'gamma': gamma}

        # No PCA
        _, acc_no_pca, auc_no_pca, y_test_val, y_proba_no_pca = evaluate_svc(
            X_train, X_test, y_train, y_test, use_pca=False
        )

        # With PCA
        _, acc_pca, auc_pca, _, y_proba_pca = evaluate_svc(
            X_train, X_test, y_train, y_test, use_pca=True, pca_variance=0.95
        )

        results.append({
            'kernel': kernel,
            'C': C,
            'gamma': gamma,
            'Accuracy_no_PCA': acc_no_pca,
            'Accuracy_PCA': acc_pca
        })

# ===== RESULTS TABLE =====
results_df = pd.DataFrame(results)
print("SVC Performance Table")
print(results_df)

# ===== BEST MODEL =====
best_idx = results_df['Accuracy_no_PCA'].idxmax() # you can also pick max of PCA
best_params = results_df.iloc[best_idx]
print("\nBest Params (No PCA)")
print(best_params)

# ===== ROC CURVE FOR BEST MODEL =====
# Using No PCA best model
best_kernel = best_params['kernel']
best_C = best_params['C']
best_gamma = best_params['gamma']

```

```
svc_best = SVC(kernel=best_kernel, C=best_C, gamma=best_gamma, probability=True)
svc_best.fit(X_train, y_train)
y_proba_best = svc_best.predict_proba(X_test)[:,-1]
fpr, tpr, _ = roc_curve(y_test, y_proba_best)
roc_auc = roc_auc_score(y_test, y_proba_best)

plt.figure(figsize=(6,5))
plt.plot(fpr, tpr, label=f'ROC curve (AUC = {roc_auc:.3f})', color='blue')
plt.plot([0,1], [0,1], 'k--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('SVC ROC Curve (Best Params)')
plt.legend(loc='lower right')
plt.show()
```



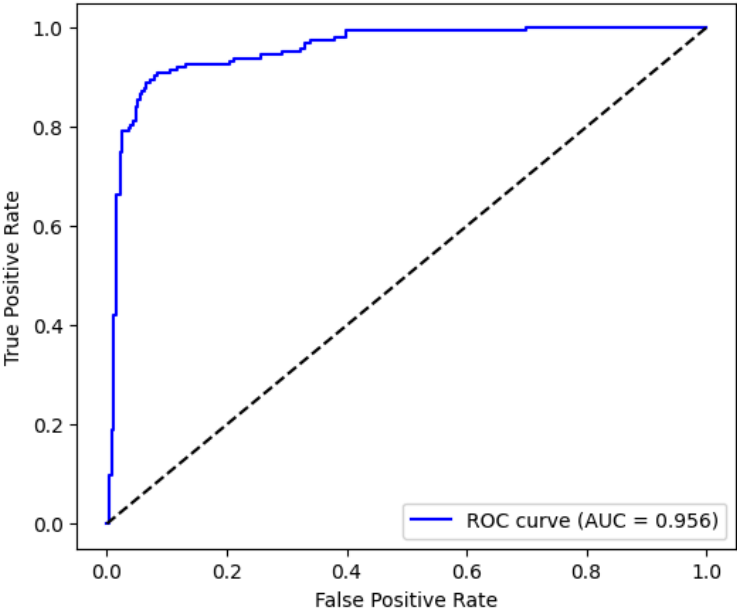
SVC Performance Table

	kernel	C	gamma	Accuracy_no_PCA	Accuracy_PCA
0	linear	0.1	scale	0.93135	0.924485
1	linear	10.0	scale	0.93135	0.924485
2	rbf	0.1	scale	0.93135	0.924485
3	rbf	10.0	scale	0.93135	0.924485

Best Params (No PCA)

kernel linear  
C 0.1  
gamma scale  
Accuracy\_no\_PCA 0.93135  
Accuracy\_PCA 0.924485  
Name: 0, dtype: object

SVC ROC Curve (Best Params)



## Naive Bayes

```
# Smoothing values to test
smoothing_values = [1e-9, 1e-8, 1e-7, 1e-6]

def evaluate_nb(X_train, X_test, y_train, y_test, smoothing, use_pca=False):
    if use_pca:
        pca = PCA(n_components=0.95)
        X_train_proc = pca.fit_transform(X_train)
        X_test_proc = pca.transform(X_test)
    else:
        X_train_proc = X_train
        X_test_proc = X_test

    model = GaussianNB(var_smoothing=smoothing)
    model.fit(X_train_proc, y_train)
    y_pred = model.predict(X_test_proc)
    y_prob = model.predict_proba(X_test_proc)[: , 1]
    acc = accuracy_score(y_test, y_pred)
    auc = roc_auc_score(y_test, y_prob)
    return acc, auc, y_prob

results = []
for s in smoothing_values:
    acc_no_pca, auc_no_pca, prob_no_pca = evaluate_nb(
        X_train, X_test, y_train, y_test, s, use_pca=False
    )
    acc_pca, auc_pca, prob_pca = evaluate_nb(
        X_train, X_test, y_train, y_test, s, use_pca=True
    )
    results.append({
        "smoothing": s,
        "Accuracy_no_PCA": acc_no_pca,
        "Accuracy_PCA": acc_pca
    })

results_df = pd.DataFrame(results)
print("Naive Bayes Performance Table")
```

```
print(results_df)

# --- Best Params (based on no PCA accuracy) ---
best_idx = results_df['Accuracy_no_PCA'].idxmax()
best_params = results_df.iloc[best_idx]
print("\nBest Smoothing (No PCA)")
print(best_params)

# --- ROC curve for best model ---
best_s = best_params['smoothing']
nb_best = GaussianNB(var_smoothing=best_s)
nb_best.fit(X_train, y_train)
y_proba_best = nb_best.predict_proba(X_test)[: , 1]
fpr, tpr, _ = roc_curve(y_test, y_proba_best)
roc_auc = roc_auc_score(y_test, y_proba_best)

plt.figure(figsize=(6,5))
plt.plot(fpr, tpr, label=f'ROC (AUC = {roc_auc:.3f})', color='purple')
plt.plot([0,1],[0,1], 'k--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('GaussianNB ROC Curve (Best Smoothing)')
plt.legend(loc='lower right')
plt.show()
```

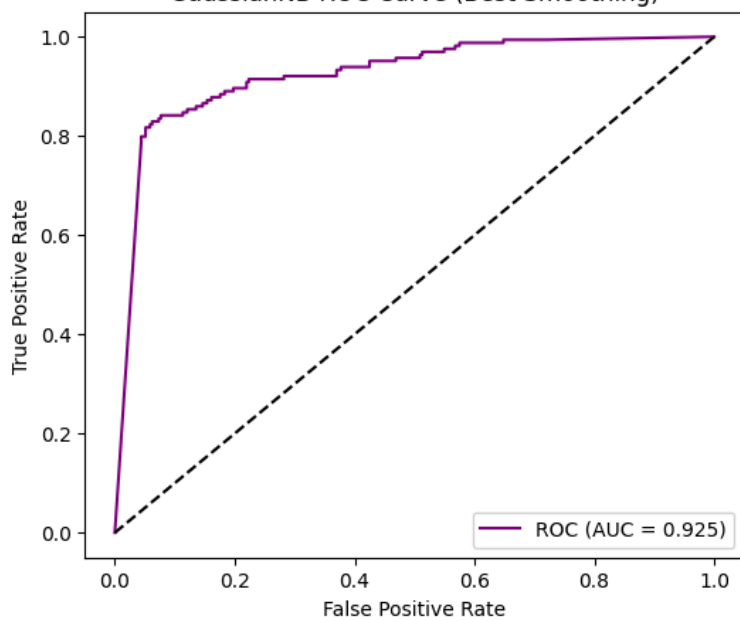
#### Naive Bayes Performance Table

	smoothing	Accuracy_no_PCA	Accuracy_PCA
0	1.000000e-09	0.720824	0.832952
1	1.000000e-08	0.723112	0.832952
2	1.000000e-07	0.725400	0.832952
3	1.000000e-06	0.725400	0.832952

#### Best Smoothing (No PCA)

smoothing 1.000000e-07  
Accuracy\_no\_PCA 7.254005e-01  
Accuracy\_PCA 8.329519e-01  
Name: 2, dtype: float64

GaussianNB ROC Curve (Best Smoothing)



#### K-Nearest Neighbours

```

# --- Hyperparam grid ---
k_values = [3, 5]
weights_options = ['uniform', 'distance']
metrics_options = ['euclidean', 'manhattan']

def evaluate_knn(X_train, X_test, y_train, y_test,
                 k, weight, metric, use_pca=False):
    if use_pca:
        pca = PCA(n_components=0.95)
        X_train_proc = pca.fit_transform(X_train)
        X_test_proc = pca.transform(X_test)
    else:
        X_train_proc, X_test_proc = X_train, X_test

    model = KNeighborsClassifier(n_neighbors=k,
                                weights=weight,
                                metric=metric)
    model.fit(X_train_proc, y_train)
    y_pred = model.predict(X_test_proc)
    y_prob = model.predict_proba(X_test_proc)[:, 1]
    acc = accuracy_score(y_test, y_pred)
    auc = roc_auc_score(y_test, y_prob)
    return acc, auc, y_prob

results = []
for k in k_values:
    for w in weights_options:
        for m in metrics_options:
            acc_no_pca, auc_no_pca, _ = evaluate_knn(
                X_train, X_test, y_train, y_test,
                k, w, m, use_pca=False
            )
            acc_pca, auc_pca, _ = evaluate_knn(
                X_train, X_test, y_train, y_test,
                k, w, m, use_pca=True
            )
            results.append({
                "k": k,

```

```

        "weights": w,
        "metric": m,
        "Accuracy_no_PCA": acc_no_pca,
        "Accuracy_PCA": acc_pca
    })

results_df = pd.DataFrame(results)
print("KNN Performance Table")
print(results_df)

# --- Pick best based on no-PCA accuracy ---
best_idx = results_df['Accuracy_no_PCA'].idxmax()
best_params = results_df.iloc[best_idx]
print("\nBest KNN Params (No PCA)")
print(best_params)

# --- ROC curve for that best combo ---
best_k = best_params['k']
best_w = best_params['weights']
best_m = best_params['metric']

knn_best = KNeighborsClassifier(n_neighbors=best_k,
                               weights=best_w,
                               metric=best_m)

knn_best.fit(X_train, y_train)
y_proba_best = knn_best.predict_proba(X_test)[:, 1]
fpr, tpr, _ = roc_curve(y_test, y_proba_best)
roc_auc = roc_auc_score(y_test, y_proba_best)

plt.figure(figsize=(6,5))
plt.plot(fpr, tpr, label=f'ROC (AUC = {roc_auc:.3f})', color='green')
plt.plot([0,1], [0,1], 'k--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('KNN ROC Curve (Best Params)')
plt.legend(loc='lower right')
plt.show()

```



## Logistic Regression

KNN Performance Table

	k	weights	metric	Accuracy_no_PCA	Accuracy_PCA
MetricRegression	0	3	uniform euclidean	0.899314	0.901602
	1	3	uniform manhattan	0.899314	0.897025
	2	3	distance euclidean	0.910755	0.908467

```
# --- Hyperparam grid ---
c_values = [0.01, 0.1, 1]
penalties = ['l2', 'l1']

def evaluate_logreg(X_train, X_test, y_train, y_test,
                   c_val, penalty, use_pca=False):
    if use_pca:
        pca = PCA(n_components=0.95)
        X_train_proc = pca.fit_transform(X_train)
        X_test_proc = pca.transform(X_test)
    else:
        X_train_proc, X_test_proc = X_train, X_test

    solver = 'saga' if penalty == 'l1' else 'lbfgs'
    model = LogisticRegression(C=c_val,
                              penalty=penalty,
                              solver=solver,
                              max_iter=5000)

    model.fit(X_train_proc, y_train)
    y_pred = model.predict(X_test_proc)
    y_prob = model.predict_proba(X_test_proc)[:, 1]
    acc = accuracy_score(y_test, y_pred)
    auc = roc_auc_score(y_test, y_prob)
    return acc, auc, y_prob

results = []
for c in c_values:
    for p in penalties:
        acc_no_pca, auc_no_pca, _ = evaluate_logreg(
            X_train, X_test, y_train, y_test,
            c, p, use_pca=False
        )
        acc_pca, auc_pca, _ = evaluate_logreg(
            X_train, X_test, y_train, y_test,
```



```

        c, p, use_pca=True
    )
    results.append({
        "C": c,
        "penalty": p,
        "Accuracy_no_PCA": acc_no_pca,
        "Accuracy_PCA": acc_pca
    })

results_df = pd.DataFrame(results)
print("Logistic Regression Performance Table")
print(results_df)

# --- Pick best by no-PCA accuracy ---
best_idx = results_df['Accuracy_no_PCA'].idxmax()
best_params = results_df.iloc[best_idx]
print("\nBest Logistic Regression Params (No PCA)")
print(best_params)

# --- ROC curve for best combo ---
best_c = best_params['C']
best_p = best_params['penalty']
solver = 'saga' if best_p == 'l1' else 'lbfgs'

log_best = LogisticRegression(C=best_c,
                              penalty=best_p,
                              solver=solver,
                              max_iter=5000)

log_best.fit(X_train, y_train)
y_proba_best = log_best.predict_proba(X_test)[:, 1]
fpr, tpr, _ = roc_curve(y_test, y_proba_best)
roc_auc = roc_auc_score(y_test, y_proba_best)

plt.figure(figsize=(6,5))
plt.plot(fpr, tpr, label=f'ROC (AUC = {roc_auc:.3f})', color='orange')
plt.plot([0,1], [0,1], 'k--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Logistic Regression ROC Curve (Best Params)')

```

```
plt.legend(loc='lower right')  
plt.show()
```

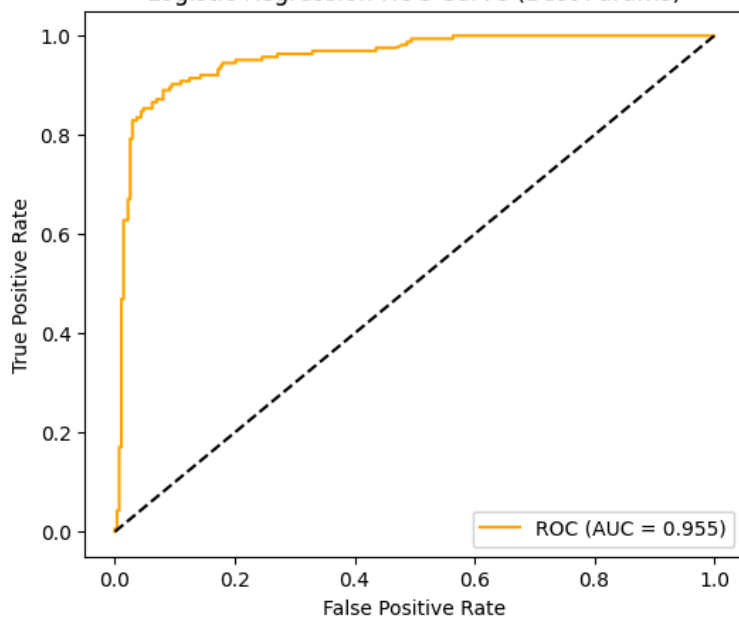
#### Logistic Regression Performance Table

	C	penalty	Accuracy_no_PCA	Accuracy_PCA
0	0.01	l2	0.908467	0.908467
1	0.01	l1	0.871854	0.894737
2	0.10	l2	0.903890	0.906178
3	0.10	l1	0.913043	0.917620
4	1.00	l2	0.910755	0.901602
5	1.00	l1	0.910755	0.906178

#### Best Logistic Regression Params (No PCA)

C 0.1  
penalty l1  
Accuracy\_no\_PCA 0.913043  
Accuracy\_PCA 0.91762  
Name: 3, dtype: object

Logistic Regression ROC Curve (Best Params)



## Decision Tree

```
# --- Hyperparams to explore ---
criteria = ["gini", "entropy"]
depths = [None, 5, 10]

def evaluate_dt(X_train, X_test, y_train, y_test,
               criterion, depth, use_pca=False):
    if use_pca:
        pca = PCA(n_components=0.95)
        X_train_proc = pca.fit_transform(X_train)
        X_test_proc = pca.transform(X_test)
    else:
        X_train_proc, X_test_proc = X_train, X_test

    model = DecisionTreeClassifier(
        criterion=criterion,
        max_depth=depth,
        random_state=42
    )
    model.fit(X_train_proc, y_train)
    y_pred = model.predict(X_test_proc)
    y_prob = model.predict_proba(X_test_proc)[:, 1]
    acc = accuracy_score(y_test, y_pred)
    auc = roc_auc_score(y_test, y_prob)
    return acc, auc, y_prob

results = []
for c in criteria:
    for d in depths:
        acc_no_pca, auc_no_pca, _ = evaluate_dt(
            X_train, X_test, y_train, y_test,
            c, d, use_pca=False
        )
        acc_pca, auc_pca, _ = evaluate_dt(
            X_train, X_test, y_train, y_test,
            c, d, use_pca=True
        )
```

```

        results.append({
            "criterion": c,
            "max_depth": d,
            "Accuracy_no_PCA": acc_no_pca,
            "Accuracy_PCA": acc_pca
        })

results_df = pd.DataFrame(results)
print("Decision Tree Performance Table")
print(results_df)

# --- Best by no-PCA accuracy ---
best_idx = results_df['Accuracy_no_PCA'].idxmax()
best_params = results_df.iloc[best_idx]
print("\nBest Decision Tree Params (No PCA)")
print(best_params)

# --- ROC curve for best combo ---
best_c = best_params['criterion']
best_d = best_params['max_depth']

# convert NaN to None, otherwise to int
if pd.isna(best_d):
    best_d = None
else:
    best_d = int(best_d)

dt_best = DecisionTreeClassifier(
    criterion=best_params['criterion'],
    max_depth=best_d,
    random_state=42
)
dt_best.fit(X_train, y_train)
y_proba_best = dt_best.predict_proba(X_test)[: , 1]
fpr, tpr, _ = roc_curve(y_test, y_proba_best)
roc_auc = roc_auc_score(y_test, y_proba_best)

plt.figure(figsize=(6,5))
plt.plot(fpr, tpr, label=f'ROC (AUC = {roc_auc:.3f})', color='brown')

```

```
plt.plot([0,1], [0,1], 'k--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Decision Tree ROC Curve (Best Params)')
plt.legend(loc='lower right')
plt.show()
```

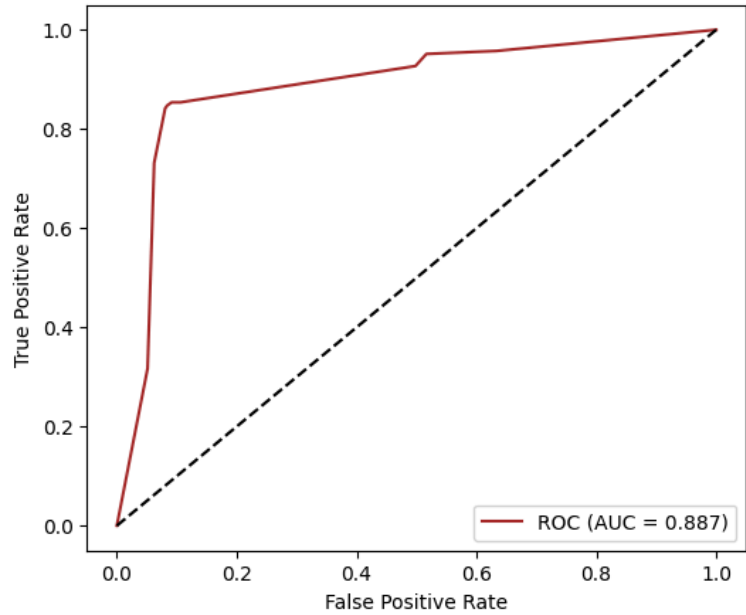
Decision Tree Performance Table

	criterion	max_depth	Accuracy_no_PCA	Accuracy_PCA
0	gini	NaN	0.887872	0.883295
1	gini	5.0	0.876430	0.860412
2	gini	10.0	0.890160	0.874142
3	entropy	NaN	0.855835	0.883295
4	entropy	5.0	0.871854	0.881007
5	entropy	10.0	0.878719	0.878719

Best Decision Tree Params (No PCA)

criterion gini  
max\_depth 10.0  
Accuracy\_no\_PCA 0.89016  
Accuracy\_PCA 0.874142  
Name: 2, dtype: object

Decision Tree ROC Curve (Best Params)



## Random Forest

```
# --- Hyperparams to explore ---
n_estimators_list = [50, 100]
max_depth_list = [None, 5, 10]

def evaluate_rf(X_train, X_test, y_train, y_test,
               n_estimators, depth, use_pca=False):
    if use_pca:
        pca = PCA(n_components=0.95)
        X_train_proc = pca.fit_transform(X_train)
        X_test_proc = pca.transform(X_test)
    else:
        X_train_proc, X_test_proc = X_train, X_test

    model = RandomForestClassifier(
        n_estimators=n_estimators,
        max_depth=depth,
        random_state=42,
        n_jobs=-1
    )
    model.fit(X_train_proc, y_train)
    y_pred = model.predict(X_test_proc)
    y_prob = model.predict_proba(X_test_proc)[: , 1]
    acc = accuracy_score(y_test, y_pred)
    auc = roc_auc_score(y_test, y_prob)
    return acc, auc, y_prob

results = []
for n in n_estimators_list:
    for d in max_depth_list:
        acc_no_pca, auc_no_pca, _ = evaluate_rf(
            X_train, X_test, y_train, y_test,
            n, d, use_pca=False
        )
        acc_pca, auc_pca, _ = evaluate_rf(
            X_train, X_test, y_train, y_test,
            n, d, use_pca=True
        )
```



```

    )
    results.append({
        "n_estimators": n,
        "max_depth": d,
        "Accuracy_no_PCA": acc_no_pca,
        "Accuracy_PCA": acc_pca
    })

results_df = pd.DataFrame(results)
print("Random Forest Performance Table")
print(results_df)

# --- Best by no-PCA accuracy ---
best_idx = results_df['Accuracy_no_PCA'].idxmax()
best_params = results_df.iloc[best_idx]
print("\nBest Random Forest Params (No PCA)")
print(best_params)

# --- Fix dtype for max_depth ---
best_depth = best_params['max_depth']
if pd.isna(best_depth):
    best_depth = None
else:
    best_depth = int(best_depth)

# --- ROC curve for best combo ---
rf_best = RandomForestClassifier(
    n_estimators=int(best_params['n_estimators']),
    max_depth=best_depth,
    random_state=42,
    n_jobs=-1
)
rf_best.fit(X_train, y_train)
y_proba_best = rf_best.predict_proba(X_test)[:, 1]
fpr, tpr, _ = roc_curve(y_test, y_proba_best)
roc_auc = roc_auc_score(y_test, y_proba_best)

plt.figure(figsize=(6,5))
plt.plot(fpr, tpr, label=f'ROC (AUC = {roc_auc:.3f})', color='darkgreen')

```

```
plt.plot([0,1], [0,1], 'k--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Random Forest ROC Curve (Best Params)')
plt.legend(loc='lower right')
plt.show()
```

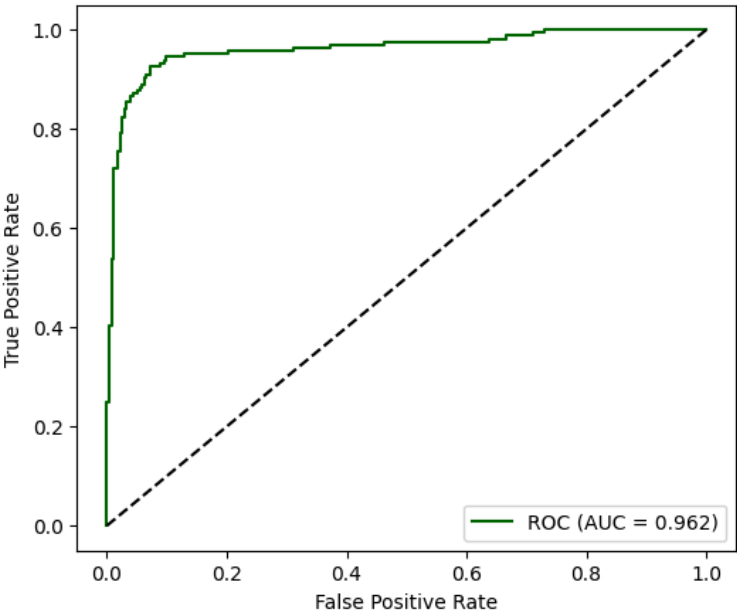
Random Forest Performance Table

	n_estimators	max_depth	Accuracy_no_PCA	Accuracy_PCA
0	50	NaN	0.922197	0.906178
1	50	5.0	0.910755	0.897025
2	50	10.0	0.924485	0.908467
3	100	NaN	0.924485	0.908467
4	100	5.0	0.908467	0.899314
5	100	10.0	0.922197	0.910755

Best Random Forest Params (No PCA)

n\_estimators 50.000000  
max\_depth 10.000000  
Accuracy\_no\_PCA 0.924485  
Accuracy\_PCA 0.908467  
Name: 2, dtype: float64

Random Forest ROC Curve (Best Params)



## AdaBoost

```
# --- Hyperparam grids ---
n_estimators_list = [50, 100]
learning_rates = [0.01, 0.1, 1.0]

def eval_adaboost(X_train, X_test, y_train, y_test,
                  n_est, lr, use_pca=False):
    if use_pca:
        pca = PCA(n_components=0.95)
        X_train_proc = pca.fit_transform(X_train)
        X_test_proc = pca.transform(X_test)
    else:
        X_train_proc, X_test_proc = X_train, X_test

    model = AdaBoostClassifier(
        n_estimators=n_est,
        learning_rate=lr,
        random_state=42
    )
    model.fit(X_train_proc, y_train)
    y_pred = model.predict(X_test_proc)
    y_prob = model.predict_proba(X_test_proc)[:, 1]
    return (
        accuracy_score(y_test, y_pred),
        roc_auc_score(y_test, y_prob)
    )

results = []
for n in n_estimators_list:
    for lr in learning_rates:
        acc_no_pca, _ = eval_adaboost(X_train, X_test, y_train, y_test,
                                       n, lr, use_pca=False)
        acc_pca, _ = eval_adaboost(X_train, X_test, y_train, y_test,
                                    n, lr, use_pca=True)
        results.append({
            "n_estimators": n,
            "learning_rate": lr,
```

```

        "Accuracy_no_PCA": acc_no_pca,
        "Accuracy_PCA": acc_pca
    })

results_df = pd.DataFrame(results)
print("AdaBoost Performance")
print(results_df)

# --- Best combo (no PCA) ---
best_idx = results_df['Accuracy_no_PCA'].idxmax()
best_row = results_df.iloc[best_idx]
print("\nBest Params (No PCA)")
print(best_row)

# --- ROC for the best ---
best_model = AdaBoostClassifier(
    n_estimators=int(best_row['n_estimators']),
    learning_rate=float(best_row['learning_rate']),
    random_state=42
)
best_model.fit(X_train, y_train)
y_best_prob = best_model.predict_proba(X_test)[:, 1]
fpr, tpr, _ = roc_curve(y_test, y_best_prob)
roc_auc = roc_auc_score(y_test, y_best_prob)

plt.figure(figsize=(6,5))
plt.plot(fpr, tpr, label=f'ROC (AUC = {roc_auc:.3f})', color='crimson')
plt.plot([0,1], [0,1], 'k--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('AdaBoost ROC Curve (Best Params)')
plt.legend(loc='lower right')
plt.show()

```

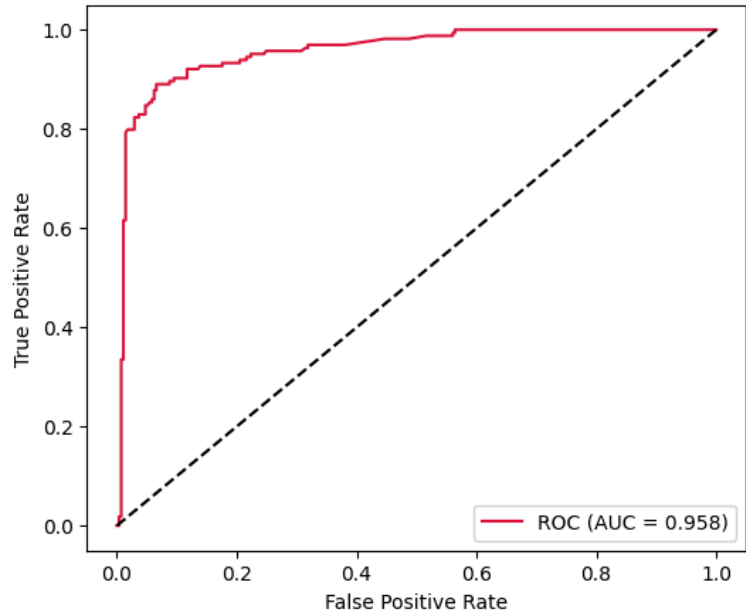
AdaBoost Performance

	n_estimators	learning_rate	Accuracy_no_PCA	Accuracy_PCA
0	50	0.01	0.821510	0.878719
1	50	0.10	0.897025	0.878719
2	50	1.00	0.908467	0.892449
3	100	0.01	0.844394	0.878719
4	100	0.10	0.899314	0.883295
5	100	1.00	0.915332	0.901602

Best Params (No PCA)

n\_estimators 100.000000  
learning\_rate 1.000000  
Accuracy\_no\_PCA 0.915332  
Accuracy\_PCA 0.901602  
Name: 5, dtype: float64

AdaBoost ROC Curve (Best Params)



## Gradient Boosting

```
# --- Hyperparam grids ---
n_estimators_list = [50, 100, 200]
learning_rates = [0.1, 0.2]

def eval_gb(X_train, X_test, y_train, y_test, n_est, lr, use_pca=False):
    if use_pca:
        pca = PCA(n_components=0.95)
        X_train_proc = pca.fit_transform(X_train)
        X_test_proc = pca.transform(X_test)
    else:
        X_train_proc, X_test_proc = X_train, X_test

    model = GradientBoostingClassifier(
        n_estimators=n_est,
        learning_rate=lr,
        random_state=42
    )
    model.fit(X_train_proc, y_train)
    y_pred = model.predict(X_test_proc)
    y_prob = model.predict_proba(X_test_proc)[:, 1]
    return accuracy_score(y_test, y_pred), roc_auc_score(y_test, y_prob)

# Collect results
results = []
for n in n_estimators_list:
    for lr in learning_rates:
        acc_no_pca, _ = eval_gb(X_train, X_test, y_train, y_test,
                                n, lr, use_pca=False)
        acc_pca, _ = eval_gb(X_train, X_test, y_train, y_test,
                              n, lr, use_pca=True)
        results.append({
            "n_estimators": n,
            "learning_rate": lr,
            "Accuracy_no_PCA": acc_no_pca,
            "Accuracy_PCA": acc_pca
        })
```

```

results_df = pd.DataFrame(results)
print("Gradient Boosting Performance")
print(results_df)

# --- Best combo (no PCA) ---
best_idx = results_df['Accuracy_no_PCA'].idxmax()
best_row = results_df.iloc[best_idx]
print("\nBest Params (No PCA)")
print(best_row)

# --- ROC curve for best model ---
best_model = GradientBoostingClassifier(
    n_estimators=int(best_row["n_estimators"]),
    learning_rate=float(best_row["learning_rate"]),
    random_state=42
)
best_model.fit(X_train, y_train)
y_best_prob = best_model.predict_proba(X_test)[:, 1]

fpr, tpr, _ = roc_curve(y_test, y_best_prob)
roc_auc = roc_auc_score(y_test, y_best_prob)

plt.figure(figsize=(6,5))
plt.plot(fpr, tpr, label=f'ROC (AUC = {roc_auc:.3f})', color='darkorange')
plt.plot([0,1], [0,1], 'k--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Gradient Boosting ROC Curve (Best Params)')
plt.legend(loc='lower right')
plt.show()

```



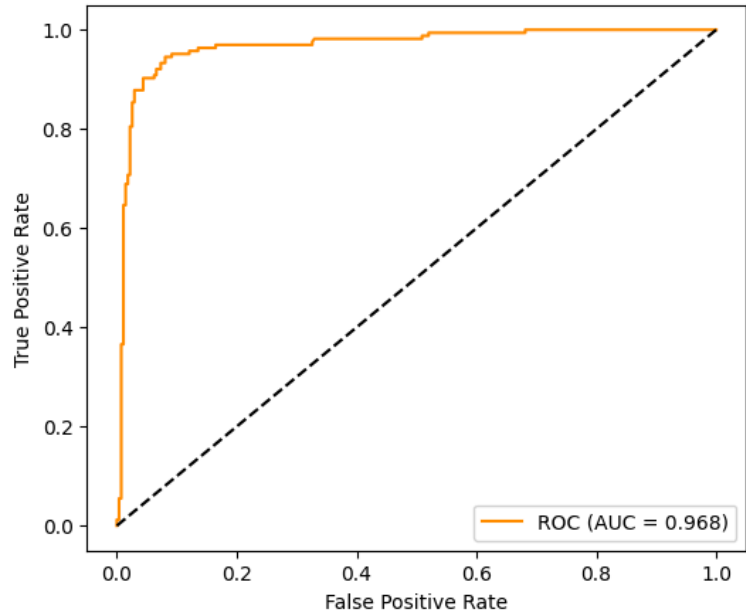
Gradient Boosting Performance

	n_estimators	learning_rate	Accuracy_no_PCA	Accuracy_PCA
0	50	0.1	0.924485	0.901602
1	50	0.2	0.922197	0.913043
2	100	0.1	0.931350	0.908467
3	100	0.2	0.929062	0.913043
4	200	0.1	0.924485	0.908467
5	200	0.2	0.931350	0.910755

Best Params (No PCA)

n\_estimators 100.000000  
learning\_rate 0.100000  
Accuracy\_no\_PCA 0.931350  
Accuracy\_PCA 0.908467  
Name: 2, dtype: float64

Gradient Boosting ROC Curve (Best Params)





```

        rows.append({
            "n_estimators": n,
            "learning_rate": lr,
            "max_depth": d,
            "Accuracy_no_PCA": acc_no,
            "Accuracy_PCA": acc_pca
        })

results_df = pd.DataFrame(rows)
print("XGBoost Performance Table")
print(results_df)

# --- Best params (No PCA) ---
best_idx = results_df['Accuracy_no_PCA'].idxmax()
best_row = results_df.iloc[best_idx]
print("\nBest Params (No PCA)")
print(best_row)

# --- ROC curve for best model ---
best_model = XGBClassifier(
    n_estimators=int(best_row["n_estimators"]),
    learning_rate=float(best_row["learning_rate"]),
    max_depth=int(best_row["max_depth"]),
    eval_metric='logloss',
    use_label_encoder=False,
    random_state=42
)
best_model.fit(X_train, y_train)
y_prob = best_model.predict_proba(X_test)[:, 1]

fpr, tpr, _ = roc_curve(y_test, y_prob)
roc_auc = roc_auc_score(y_test, y_prob)

plt.figure(figsize=(6,5))
plt.plot(fpr, tpr, label=f'ROC (AUC={roc_auc:.3f})', color='crimson')
plt.plot([0,1], [0,1], 'k--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('XGBoost ROC Curve (Best Params)')

```

```
plt.legend(loc='lower right')  
plt.show()
```



```
/usr/local/lib/python3.12/dist-packages/xgboost/training.py:183: UserWarning: [15:42:51] WARNING: /workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

```
bst.update(dtrain, iteration=i, fobj=obj)
/usr/local/lib/python3.12/dist-packages/xgboost/training.py:183: UserWarning: [15:42:51] WARNING: /workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

```
bst.update(dtrain, iteration=i, fobj=obj)
/usr/local/lib/python3.12/dist-packages/xgboost/training.py:183: UserWarning: [15:42:51] WARNING: /workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

```
bst.update(dtrain, iteration=i, fobj=obj)
/usr/local/lib/python3.12/dist-packages/xgboost/training.py:183: UserWarning: [15:42:52] WARNING: /workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

### Stacking (base learners + meta-learner)

```
bst.update(dtrain, iteration=i, fobj=obj)
/usr/local/lib/python3.12/dist-packages/xgboost/training.py:183: UserWarning: [15:42:53] WARNING: /workspace/src/learner.cc:738:
```

```
# --- Base / final models ---
svm = SVC(kernel='linear', probability=True, random_state=42)
nb = GaussianNB()
dt = DecisionTreeClassifier(random_state=42)
knn = KNeighborsClassifier()

log_reg = LogisticRegression(max_iter=1000, random_state=42)
rf = RandomForestClassifier(n_estimators=100, random_state=42)

# 3 stack varieties
stacks = {
    "SVM+NB+DT → LR": StackingClassifier(
        estimators=[('svm', svm), ('nb', nb), ('dt', dt)],
        final_estimator=log_reg, passthrough=False, n_jobs=-1
    ),
    "SVM+NB+DT → RF": StackingClassifier(
        estimators=[('svm', svm), ('nb', nb), ('dt', dt)],
        final_estimator=rf, passthrough=False, n_jobs=-1
    ),
    "SVM+DT+KNN → LR": StackingClassifier(
        estimators=[('svm', svm), ('dt', dt), ('knn', knn)],
        final_estimator=log_reg, passthrough=False, n_jobs=-1
    )
}
```

```

}

def eval_stack(model, use_pca=False):
    if use_pca:
        pca = PCA(n_components=0.95)
        Xtr = pca.fit_transform(X_train)
        Xte = pca.transform(X_test)
    else:
        Xtr, Xte = X_train, X_test
    model.fit(Xtr, y_train)
    y_pred = model.predict(Xte)
    y_prob = model.predict_proba(Xte)[:, 1]
    return accuracy_score(y_test, y_pred), roc_auc_score(y_test, y_prob), y_prob

results = []
roc_curves = {}

for name, model in stacks.items():
    acc_no, auc_no, prob_no = eval_stack(model, use_pca=False)
    acc_pca, auc_pca, prob_pca = eval_stack(model, use_pca=True)
    results.append({
        "Model": name,
        "Accuracy_no_PCA": acc_no,
        "Accuracy_PCA": acc_pca,
        "AUC_no_PCA": auc_no,
        "AUC_PCA": auc_pca
    })

```

# Hyperparameter Tuning Tables

**Table 1: PCA Summary**

Setting	Variance Target	Explained Variance (%)	
With PCA	95%	95.54	Captures 95% of total variance while reducing dimensionality

**Table 2: SVM - Hyperparameter Tuning Results**

Kernel	C	Gamma	No-PCA	With-PCA
linear	0.1	scale	0.93135	0.924485
linear	10.0	scale	0.93135	0.924485
rbf	0.1	scale	0.93135	0.924485
rbf	10.0	scale	0.93135	0.924485

**Table 3: Naive Bayes - Smoothing Choices**

Smoothing	No-PCA	With-PCA
1e-09	0.720824	0.832952
1e-08	0.723112	0.832952
1e-07	0.725400	0.832952
1e-06	0.725400	0.832952

**Table 4: K-Nearest Neighbors (KNN)**

k	Weights	Metric	No-PCA	With-PCA
3	uniform	euclidean	0.899314	0.901602
3	distance	euclidean	0.910755	0.908467
5	distance	euclidean	0.908467	0.910755

**Table 5: Logistic Regression**

C	Penalty	No-PCA	With-PCA
0.01	l2	0.908467	0.908467
0.10	l1	0.913043	0.917620
1.00	l2	0.910755	0.901602

**Table 6: Decision Tree**

Criterion	Max Depth	No-PCA	With-PCA
gini	10	0.890160	0.874142
entropy	10	0.878719	0.878719

**Table 7: Random Forest**

N Estimators	Max Depth	No-PCA	With-PCA
50	10	0.924485	0.908467
100	10	0.922197	0.910755



**Table 8: AdaBoost**

N Estimators	Learning Rate	No-PCA	With-PCA
50	1.00	0.908467	0.892449
100	1.00	0.915332	0.901602

**Table 9: Gradient Boosting**

N Estimators	Learning Rate	No-PCA	With-PCA
100	0.1	0.931350	0.908467
100	0.2	0.929062	0.913043

**Table 10: XGBoost**

N Estimators	Learning Rate	Max Depth	No-PCA	With-PCA
100	0.1	7	0.940503	0.910755

**Table 11: Stacked Models**

Model	No-PCA	With-PCA
SVM+NB+DT $\rightarrow$ LR	0.917620	0.913043
SVM+DT+KNN $\rightarrow$ LR	0.919908	0.919908

## Observations

- Naive Bayes and KNN benefited most from PCA. Ensemble models (RF, GB, XGB) improved moderately.
- PCA reduced variance across folds, indicating more stable results.
- PCA helped reduce overfitting in simpler models like NB and KNN.
- Linear models (SVM, Logistic Regression) showed minor performance change.
- Stacking remained robust to dimensionality reduction.

## Learning Outcomes

- Learned to perform hyperparameter tuning for multiple ML classifiers.
- Applied PCA and studied its impact on model accuracy and variance.
- Identified which models benefit from PCA (e.g., NB, KNN) and which remain robust (e.g., ensembles).
- Understood that stacking ensembles maintain performance even after dimensionality reduction.

**GitHub Repository:** [:github.com/Thamizhmathibharathi/project/tree/main/assignment6](https://github.com/Thamizhmathibharathi/project/tree/main/assignment6)

## References

1. A. Lubis, P. Sihombing, and E. Nababan, “Analysis of accuracy improvement in k-nearest neighbor using principal component analysis (PCA),” *Journal of Physics: Conference Series*, vol. 1566, p. 012062, June 2020.
2. M. Salam, A. Azar, M. Elgendy, and K. Fouad, “The effect of different dimensionality reduction techniques on machine learning overfitting problem,” *International Journal of Advanced Computer Science and Applications*, vol. 12, no. 4, pp. 641–655, 2021.