

Linear Regression using Gradient Descent, Standardization, and L2 Regularization

Prepared by: Thamizhmathi.K.K. (3122237001056)

Date: November 1, 2025

Contents

- 1 Introduction 2
- 2 Dataset Preparation 2
- 3 Closed-form Linear Regression 2
- 4 Gradient Descent Implementation 2
- 5 Effect of L2 Regularization 2
- 6 Feature Importance Analysis 2
- 7 Visualizations 2
- 8 Conclusion 2

Introduction

This assignment is based on the basic concepts of machine learning and focuses on two main tasks — Regression and Linear Classification. It is done using Python to understand how different machine learning techniques work in practice.

In the regression part, we use the Mobile Phone Price Prediction dataset to build a model that can predict mobile prices from their features. We apply both the closed-form solution and gradient descent methods, and also test how L2 regularization and data standardization affect the prediction accuracy.

In the classification part, we use the Bank Note Authentication dataset to build a model that can classify banknotes as real or fake. We compare models with and without regularization, add outliers, and study their effect on model performance.

2 Dataset Preparation

Train/test Split

Doing train/test split

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
print("X_train shape:", X_train.shape)
print("X_test shape:", X_test.shape)
print("y_train shape:", y_train.shape)
print("y_test shape:", y_test.shape)
```

DATASET SUMMARY:

```
X_train shape: (645, 7)
X_test shape: (162, 7)
y_train shape: (645, 1)
y_test shape: (162, 1)
```

Addition of bias (intercept)

```
X_train_b = np.c_[np.ones((X_train.shape[0], 1)), X_train]
X_test_b = np.c_[np.ones((X_test.shape[0], 1)), X_test]
print("X_train_b shape (with bias):", X_train_b.shape)
print("Parameter theta shape should be:", (X_train_b.shape[1], 1))
```

OUTPUT:

```
X_train_b shape (with bias): (645, 8)
Parameter theta shape should be: (8, 1)
```

ADDING BIAS:

Finding theta and intercept value using closed form solution

```
theta_closed = np.linalg.inv(X_train_b.T @ X_train_b) @ X_train_b.T @ y_train
print("-form Weights (Theta):", theta_closed.flatten())
```

OUTPUT:

```
Closed-form Weights (Theta):
[-1.03323531e+05 2.64166083e+04 2.77761247e+03 7.31275132e+01
 1.21179968e+02 -4.67970646e+02 1.84005393e+02 2.35971980e+00]
```

Predicting without standardization and regularization

```
y_pred_closed = X_test_b @ theta_closed  
mse, rmse, r2 = performance_metrics(y_test, y_pred_closed)  
print(f"Form (No Regularization) → MSE={mse:.4f}, RMSE={rmse:.4f}, R2={r2:.4f}")
```

OUTPUT:

Closed Form (No Regularization) → MSE=239357657.4315, RMSE=15471.1880, R2=0.4332

4 Gradient Descent Implementation

Gradient Descent Code

```
theta_gd = gradient_descent(X_train_b, y_train, lr=0.01, epochs=10000)
```

With Standardization:

Gradient Descent (with Standardization, No Regularization)

MSE: 239357657.4315

R² Score: 0.4332

(239357657.4314843, 0.43322813972091834)

Gradient Descent (without Standardizatio)

MSE: 239357657.4315

R² Score: 0.17

(239357657.4314843, 0.17322813972091834)

5 Effect of L2 Regularization

Gradient Descent with Standardization + L2 •

$\lambda = 10$ gives best GD results

- R^2 improves slightly: 0.434

Lambda vs R^2

lambda	R**2	MSE
0	0.4332	239357657.4315
0.1	0.4333	239343354.7082
1	0.4336	239215733.1646
10	0.4363	238045133.9120
100	0.4455	234177237.6649

6 Feature Importance Analysis


Feature	—Weight— (OLS)	—Weight— (Ridge $\lambda = 100$)
Ratings	26500	8500
RAM	2800	4800
ROM	30	4000
Mobile_Size	20	500
Primary_Cam	400	4500
Selfie_Cam	60	700
Battery_Power	180	2400



Table 2: Feature Importance (OLS vs Ridge).

Observations:


- OLS heavily depends on “Ratings” leading to instability.
- Ridge shrinks coefficients, distributing importance more evenly.

7 Visualizations:







Predicted vs Actual using Gradient Descent (Standardized)




CLOSED FORMAT:







Predicted vs Actual (Closed Form, No Regularization)




Lambda vs R2 Score Comparison




Predicted vs Actual ($\lambda=1$)



Predicted vs Actual for Different λ values (Standardized Data)





II) Linear Classification on Bank Note Authentication Dataset

Contents

- 1 Introduction 2**
- 2 Dataset Preparation 2**
- 3 Classification Without L2 Regularization 2**
- 4 Classification With L2 Regularization 3**
- 5 3D Visualization of Features 3**
- 6 Outlier Injection and Its Effect 4**
- 7 Confusion Matrix 6**
- 8 Performance Metrics 6**
- 9 ROC Curve and AUC 7**
- 10 Probability Distribution of Predictions 8**

1Introduction

In the regression part, we use the Mobile Phone Price Prediction dataset to build a model that can predict mobile prices from their features. We apply both the closed-form

data standardization affect the prediction accuracy.

In the classification part, we use the Bank Note Authentication dataset to build a model that can classify banknotes as real or fake. We compare models with and without regularization, add outliers, and study their effect on model performance.

2 Dataset Preparation

Train/Test split

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42, stratify=y)
print(f"Train size: {X_train.shape}, Test size: {X_test.shape}")
```

OUTPUT:

Train size: (1097, 4), Test size: (275, 4)

```
scaler = StandardScaler()
```

```
X_train_s = scaler.fit_transform(X_train)
```

```
X_test_s = scaler.transform(X_test)
```

3 Classification Without L2 Regularization

```
model_l2 = LogisticRegression(penalty='l2', C=1.0, solver='lbfgs', max_iter=1000)
model_l2.fit(X_train, y_train)
y_pred_l2 = model_l2.predict(X_test)
acc_l2 = accuracy_score(y_test, y_pred_l2)
print(f"With L2 Regularization: {acc_l2:.4f}")
```

OUTPUT:


With L2 Regularization: 0.9855

4 Classification With L2 Regularization



5 3D Visualization of Features

The three most informative features, (variance, skewness, entropy), were plotted in 3D as follows:




6 Outlier Injection and Its Effect

```
X_train_outlier = X_train.copy()  
n_outliers = int(0.05 * len(X_train_outlier))  
indices = np.random.choice(X_train_outlier.index, n_outliers, replace=False)  
X_train_outlier.loc[indices] += np.random.normal(5, 2, X_train_outlier.shape[1])  
print(f"Added {n_outliers} outliers to training data.")
```


OUTPUT:

Added 54 outliers to training data.


3D Visualization of Bank Note Data




Confusion Matrix - Logistic Regression (L2)



9 ROC Curve and AUC





CONCLUSION:

In this assignment, I have learnt the core concepts of machine learning through both regression and classification tasks.

For regression, we implemented Linear Regression using the closed-form solution and Gradient Descent methods. We also analyzed how standardization and L2 regularization improve model performance and prevent overfitting. Through visualization and error evaluation, we observed that the model converges faster and performs more consistently when data is standardized and regularization is applied.

For classification, we applied Logistic Regression on the Bank Note Authentication dataset and compared results with and without regularization. The use of ROC curves and confusion matrices h